

Generic pattern unification: a categorical approach

Abstract—We provide a generic categorical setting for Miller’s pattern unification. The syntax with metavariables is generated by a free monad applied to finite coproducts of representable functors; the most general unifier is computed as a coequaliser in the Kleisli category restricted to such coproducts. Our setting handles simply-typed second-order syntax, linear syntax, or (intrinsic) polymorphic syntax such as system F.

Index Terms—Unification, Category theory, Syntax

I. INTRODUCTION

Unification consists in finding a *unifier* of two terms t, u , that is a (metavariable) substitution σ such that $t[\sigma] = u[\sigma]$. Unification algorithms try to compute a most general unifier σ , in the sense that given any other unifier δ , there exists a unique δ' such that $\delta = \sigma[\delta']$. First-order unification [19] is used in ML-style type inference systems and logic programming languages such as Prolog. For more advanced type systems, where variable binding is crucially involved, one needs second-order unification [13], which is undecidable [10]. However, Miller [15] identified a decidable fragment: in so-called *pattern unification*, metavariables are allowed to take distinct variables as arguments. In this situation, we can write an algorithm that either fails in case there is no unifier, either computes the most general unifier.

Recent results in type inference, Dunfield-Krishnaswami [5], or Jinxu et al [23], include very large proofs: the former comes with a 190 page appendix, and the latter comes with a Coq proof is many thousands of lines long -- and both of these results are for tiny kernel calculi. If we ever hope to extend this kind of result to full programming languages like Haskell or OCaml, we must raise the abstraction level of these proofs, so that they are no longer linear (with a large constant) in the size of the calculus. A close examination of these proofs shows that a large part of the problem is that the type inference algorithms make use of unification, and the correctness proofs for type inference end up essentially re-establishing the entire theory of unification for each algorithm. The reason they do this is because algorithmic typing rules essentially give a first-order functional program with no abstractions over (for example) a signature for the unification algorithm to be defined over, or any axiomatic statement of the invariants the algorithmic typing rules had to maintain.

The present work is a first step towards a general solution to this problem. Our generic unification algorithm is parameterised by an abstract notion of signature, covering simply-typed second-order syntax, linear syntax, or (intrinsic) polymorphic syntax such as system F. We focused on Miller pattern unification, a decidable fragment of higher-order unification where metavariables can only take distinct variables as arguments. This is already a step beyond the above-cited

works [23], [5] that use plain first-order unification. Moreover, this is necessary for types with binders (e.g., fixed-point operators like $\mu a. A[a]$) as well as for rich type systems like dependent types.

As an introduction, we start by presenting pattern unification in the case of pure λ -calculus in Section §I-A. In Section §I-B, we then present the generic algorithm summarised in Figure 1, instantiated for a syntax specified by a *binding signature*. Finally, in Section §I-C, we motivate our general setting and provide categorical semantics of the algorithm, by revisiting pure λ -calculus.

Related work

First-order unification has been explained from a lattice-theoretic point of view by Plotkin [3], and later categorically analysed in [20], [9], [1, Section 9.7] as coequalisers. However, there is little work on understanding pattern unification algebraically, with the notable exception of [22], working with normalised terms of simply-typed λ -calculus. The present paper can be thought of as a generalisation of their work.

Although our notion of signature has a broader scope since we are not specifically focusing on syntax where variables can be substituted, our work is closer in spirit to the presheaf approach [7] than to the nominal sets’ [8] in that everything is explicitly scoped: terms come with their support, metavariables always appear with their scope of allowed variables.

Nominal unification [21] is an alternative to pattern unification where metavariables are not supplied with the list of allowed variables. Instead, substitution can capture variables. Nominal unification explicitly deals with α -equivalence as an external relation on the syntax, and as a consequence deals with freshness problems in addition to unification problems.

A. An example: pure λ -calculus.

Consider the syntax of pure λ -calculus extended with metavariables satisfying the pattern restriction, encoded with De Bruijn levels, rather than De Bruijn indices [4]. More formally, the syntax is inductively generated by the following inductive rules, where Γ is a metavariable context $M_1 : m_1, \dots, M_p : m_p$ specifying a metavariable symbol M_i together with its number of arguments m_i .

$$\frac{1 \leq i \leq n}{\Gamma; n \vdash v_i} \quad \frac{\Gamma; n \vdash t \quad \Gamma; n \vdash u}{\Gamma; n \vdash t u} \quad \frac{\Gamma; n+1 \vdash t}{\Gamma; n \vdash \lambda t}$$

$$\frac{M : m \in \Gamma \quad 1 \leq i_1, \dots, i_m \leq n \quad i_1, \dots, i_m \text{ distinct}}{\Gamma; n \vdash M(v_{i_1}, \dots, v_{i_m})}$$

Note that the De Bruijn level convention means that the variable bound in $\Gamma; n \vdash \lambda t$ is $n+1$.

A *metavariable substitution* $\sigma : \Gamma \rightarrow \Gamma'$ assigns to each declaration $M : m$ in Γ a term $\Gamma'; m \vdash \sigma_M$. This assignment

extends (through a recursive definition) to any term $\Gamma; n \vdash t$, yielding a term $\Gamma'; n \vdash t[\sigma]$. The base case is

$$M(x_1, \dots, x_n)[\sigma] = \sigma_M[v_i \mapsto x_i], \quad (1)$$

where $-[v_i \mapsto x_i]$ is variable renaming. Composition of substitutions $\sigma : \Gamma_1 \rightarrow \Gamma_2$ and $\sigma' : \Gamma_2 \rightarrow \Gamma_3$ is then defined as $(\sigma[\sigma'])_M = \sigma_M[\sigma']$.

A *unifier* of two terms $\Gamma; n \vdash t, u$ is a substitution $\sigma : \Gamma \rightarrow \Gamma'$ such that $t[\sigma] = u[\sigma]$. A *most general unifier* of t and u is a unifier $\sigma : \Gamma \rightarrow \Gamma'$ that uniquely factors any other unifier $\delta : \Gamma \rightarrow \Delta$, in the sense that there exists a unique $\delta' : \Gamma' \rightarrow \Delta$ such that $\delta = \sigma[\delta']$. We denote this situation by $\Gamma \vdash t = u \Rightarrow \sigma \dashv \Gamma'$, leaving the variable context n implicit. Intuitively, the symbol \Rightarrow separates the input and the output of the unification algorithm, which either returns a most general unifier, either fails when there is no unifier at all (for example, when unifying $t_1 t_2$ with λu). To handle the latter case, we add a formal error metavariable context \perp in which the only term (in any variable context) is a formal error term $!$, inducing a unique substitution $! : \Gamma \rightarrow \perp$, satisfying $t[!] = !$ for any term t . For example, we have $\Gamma \vdash t_1 t_2 = \lambda u \Rightarrow ! \dashv \perp$. Formally, we will interpret \perp as a freely added terminal object in Section §III.

We generalise the notation (and thus the input of the unification algorithm) to lists of terms $\vec{t} = (t_1, \dots, t_n)$ and $\vec{u} = (u_1, \dots, u_n)$ such that $\Gamma; n_i \vdash t_i, u_i$. Then, $\Gamma \vdash \vec{t} = \vec{u} \Rightarrow \sigma \dashv \Gamma'$ means that σ unifies each pair (t_i, u_i) and is the most general one, in the sense that it uniquely factors any other substitution that unifies each pair (t_i, u_i) . As a consequence, we get the following *congruence* rule for application.

$$\frac{\Gamma \vdash t_1, t_2 = u_1, u_2 \Rightarrow \sigma \dashv \Delta}{\Gamma \vdash t_1 t_2 = u_1 u_2 \Rightarrow \sigma \dashv \Delta}$$

Unifying a list of term pairs $\vec{t}_1, \vec{t}_2 = u_1, \vec{u}_2$ can be performed sequentially by first computing the most general unifier σ_1 of (t_1, u_1) , then applying the substitution to (\vec{t}_2, \vec{u}_2) , and finally computing the most general unifier of the resulting list of term pairs: this is precisely the rule U-SPLIT in Figure 1.

Thanks to this rule, we can focus on unification of a single term pair. The idea here is to recursively inspect the structure of the given terms, until reaching a metavariable application $M(x_1, \dots, x_n)$ at top level on either hand side of $\Gamma, M : n \vdash t, u$. Assume by symmetry $t = M(x_1, \dots, x_n)$, then three mutually exclusive situations must be considered:

- 1) M appears deeply in u
- 2) M appears in u at top level, i.e., $u = M(y_1, \dots, y_n)$;
- 3) M does not appear in u ;

In the first case, there is no unifier because the size of both hand sides can never match after substitution. This justifies the rule

$$\frac{u \neq M(\dots) \quad u|_\Gamma \neq \dots}{\Gamma, M : m \vdash M(\vec{x}) = u \Rightarrow ! \dashv \perp}$$

where $u|_\Gamma \neq \dots$ means that u does not restrict to the smaller metavariable context Γ , and thus that M does appear in u .

In the second case, we are unifying $M(\vec{x})$ with $M(\vec{y})$. The most general unifier substitutes M with $M'(z_1, \dots, z_p)$, where

z_1, \dots, z_p is the family of common positions i such that $x_i = y_i$. We denote¹ such a situation by $n \vdash \vec{x} = \vec{y} \Rightarrow \vec{z} \dashv p$. We therefore get the rule

$$\frac{m \vdash \vec{x} = \vec{y} \Rightarrow \vec{z} \dashv p}{\Gamma, M : m \vdash M(\vec{x}) = M(\vec{y}) \Rightarrow M \mapsto M'(\vec{z}) \dashv \Gamma, M' : p} \quad (2)$$

The last case is unification of $M(\vec{x})$ with some u such that M does not appear in u , i.e., u restricts to the smaller metavariable context Γ . We denote such a situation by $u|_\Gamma = \underline{u}'$, where u' is essentially u but considered in the smaller metavariable context Γ . In this case, the algorithm enters a *pruning phase*. To give an example, when unifying an application $t u$ with a metavariable $M(x_1, \dots, x_n)$ which does not occur in t, u , two fresh metavariables M_1 and M_2 are created. Then, t is unified with $M_1(x_1, \dots, x_n)$, outputting a unifier σ_1 , and $u[\sigma_1]$ is unified with $M_2(x_1, \dots, x_n)$, outputting a unifier σ_2 . Eventually, M is replaced with $(M_1(\vec{x}) M_2(\vec{x}))[\sigma_2]$, where σ_2 is the output unifier. We call it *pruning* because unifying u' and $M(\vec{x})$ when M does not occur in u' consists in removing all *outbound* variables in u' , i.e., those that are not among the arguments x_1, \dots, x_n of the metavariable, by producing a substitution that restricts the arities of the metavariables occurring in u' .

Let us introduce a specific notation for this phase: $\Gamma \vdash u' :> M(\vec{x}) \Rightarrow w; \sigma \dashv \Delta$ means that σ is the output pruning substitution, and w is essentially $u'[\sigma][x_i \mapsto v_i]$, the term that the metavariable M ought to be substituted with. Note that the metavariable symbol M is fresh in this notation: it appears neither in Γ nor u' , and is not in the domain of σ .

The output σ, w defines a substitution $(\Gamma, M : n) \rightarrow \Delta$ which can be characterised as the most general unifier of u' and $M(\vec{x})$. We thus have the rule

$$\frac{u|_\Gamma = \underline{u}' \quad \Gamma \vdash u' :> M(\vec{x}) \Rightarrow w; \sigma \dashv \Delta}{\Gamma, M : n \vdash M(\vec{x}) = u \Rightarrow \sigma, M \mapsto w \dashv \Delta} \quad (3)$$

Note that M is indeed substituted by w , as hinted above. As before, we generalise the pruning phase to handle lists $\vec{u} = (u_1, \dots, u_n)$ of terms such that $\Gamma; C_i \vdash u_i$, and lists of pruning patterns $(\vec{x}_1, \dots, \vec{x}_n)$ where each \vec{x}_i is a choice of distinct variables in C_i . Then, $\Gamma \vdash \vec{u} :> M_1(\vec{x}_1), \dots, M_n(\vec{x}_n) \Rightarrow \vec{w}; \sigma \dashv \Delta$ means that σ is the common pruning substitution, and w_i is essentially $u_i[\sigma][x_{i,j} \mapsto v_j]$, which M_i ought to be substituted with. Again, (σ, \vec{w}) define a substitution from $\Gamma, M_1 : |\vec{x}_1|, \dots, M_n : |\vec{x}_n|$ to Δ which can be characterised as the most general unifier of \vec{u} and $M_1(\vec{x}_1), \dots, M_n(\vec{x}_n)$.

We can then handle application as follows.

$$\frac{\Gamma \vdash t, u :> M_1(\vec{x}), M_2(\vec{x}) \Rightarrow w_1, w_2; \sigma \dashv \Delta}{\Gamma \vdash t u :> M(\vec{x}) \Rightarrow w_1 w_2; \sigma \dashv \Delta} \quad (4)$$

As for unification (rule U-SPLIT), pruning can be done sequentially, as in the rule P-SPLIT in Figure 1. The usage of $+$ as a separator will be formally justified later in Remark 5: it intuitively enforces that the metavariables used in \vec{f}_2 are

¹The similarity with the above introduced notation is no coincidence: as we will see, both are coequalisers.

distinct from the metavariable used in f_1 . Thanks to this sequential rule, we can focus on pruning a single term. The variable case is straightforward.

$$\frac{y = x_i}{\Gamma \vdash y :> M(\vec{x}) \Rightarrow v_i; 1_\Gamma \dashv \Gamma} \quad \frac{y \notin \vec{x}}{\Gamma \vdash y :> M(\vec{x}) \Rightarrow !; ! \dashv \perp} \quad (5)$$

In λ -abstraction, the bound variable $n + 1$ need not be pruned: we extend the list of allowed variables accordingly.

$$\frac{\Gamma \vdash t :> M'(\vec{x}, n + 1) \Rightarrow w; \sigma \dashv \Delta}{\Gamma \vdash \lambda t :> M(\vec{x}) \Rightarrow \lambda w; \sigma \dashv \Delta} \quad (6)$$

The remaining case consists in unifying $N(\vec{x})$ and $M(\vec{y})$, or equivalently, pruning a metavariable $N(x_1, \dots, x_n)$, whose arity must then be restricted to those positions in y_1, \dots, y_m . To be more precise, consider the family z_1, \dots, z_p of common values in x_1, \dots, x_n and y_1, \dots, y_m , so that $z_i = x_{l_i} = y_{r_i}$ for some lists (l_1, \dots, l_p) and (r_1, \dots, r_p) of distinct elements of $\{0, \dots, n-1\}$ and $\{0, \dots, m-1\}$ respectively. We denote² such a situation by $n \vdash \vec{x} :> \vec{y} \Rightarrow \vec{l}; \vec{r} \dashv p$. Then, the metavariable N is substituted with $N'(\vec{r})$ for some new metavariable N' of arity p , while the metavariable M is replaced with $N'(\vec{l})$:

$$\frac{n \vdash \vec{x} :> \vec{y} \Rightarrow \vec{l}; \vec{r} \dashv p}{\Gamma, N : n \vdash N(\vec{x}) :> M(\vec{y}) \Rightarrow N'(\vec{l}); N \mapsto N'(\vec{r}) \dashv \Gamma, N' : p} \quad (7)$$

This ends our description of the unification algorithm, in the specific case of pure λ -calculus. The goal of this paper is to generalise it, by parameterising the algorithm by a signature specifying a syntax.

B. First generalisation: parameterisation by a binding signature

As a first step, let us parameterise the unification algorithm by a binding signature [17]. A syntax is then specified by a set of symbols O together with a list of natural numbers $\vec{\alpha}_o$ for each $o \in O$ specifying the number of arguments (the size of the list) and the number of bound variables in each argument. For example, pure λ -calculus is specified by $O = \{app, lam\}$ with $\vec{\alpha}_{app} = (0, 0)$ and $\vec{\alpha}_{lam} = (1)$. The unification algorithm described in the previous section straightforwardly generalises to any syntax specified by a binding signature. Figure 1 summarises the generic algorithm that we will later interpret in a more general setting, where metavariable arguments are morphisms in a category. Since nothing enforces them to be lists³, the vector notation is dropped for these arguments in the figure, but we still use it in the following specialisation to syntax specified by a binding signature.

In the rule U-RIGRIG, the expression $o(\vec{t})$ can be an operation or a variable, in which case \vec{t} is the empty list. If o is an operation, the exact nature of \vec{t} depends on the arity $(\alpha_1, \dots, \alpha_p)$ of o : then \vec{t} is a list of terms of size p

²Again, the similarity with the pruning notation is no coincidence: as we will see, both are pullbacks.

³See Section §VIII-B for an example where arguments are sets.

and $\Gamma; n + \alpha_i \vdash t_i$ for each $i \in \{1, \dots, p\}$, where n is the variable context of $o(\vec{t})$. The rigid case in the pruning phase consists in two rules P-RIG and P-FAIL. Both are concerned with pruning a non metavariable term $\Gamma; n \vdash o(\vec{t})$. In the variable case, these two rules instantiate to (5). More precisely, if o is a variable in n , the side condition $o = \vec{x} \cdot o'$ means that $o = x_{o'+1}$. On the other hand, if o is an operation, then $\vec{x} \cdot o$ is defined as o and thus the rule P-RIG always applies with $o' = o$. If $\alpha_o = (n_1, \dots, n_p)$, then the notation $\mathcal{L}^+ \vec{x}^o$ essentially unfolds to a list of the same size as α_o and whose i^{th} element is $M_i(\vec{x}, |C|, \dots, |C| + n_i - 1)$. For example, for pure λ -calculus, $\mathcal{L}^+ \vec{x}^o = M_1(\vec{x}), M_2(\vec{x})$ in the application case, and $\mathcal{L}^+ \vec{x}^o = M_1(\vec{x}, |C|)$ in the abstraction case, thus recovering the rules (4) and (6).

Note that the premises of the rules U-FLEXFLEX and P-FLEX are not explicitly defined in figure 1, although for a syntax specified by a binding signature, they have the same meaning as in the previous section. In fact, the generic algorithm works in a more general setting, as we are going to explain in the next section, so that they need to be customised for each specific situation.

C. Categorification

In this section, we define the syntax of pure λ -calculus from a categorical point of view in order to motivate our general categorical setting. We then explain the semantics of the generic unification algorithm summarised in Figure 1.

Consider the category of functors $[\mathbb{F}_m, \text{Set}]$ from \mathbb{F}_m , the category of finite cardinals and injections between them, to the category of sets. A functor $X : \mathbb{F}_m \rightarrow \text{Set}$ can be thought of as assigning to each natural number n a set X_n of expressions with free variables taken in the set $\underline{n} = \{v_1, \dots, v_n\}$. The action on morphisms of \mathbb{F}_m means that these expressions support injective renamings. Note that this structure is no more than what is needed to substitute a metavariable in Equation (1). Let us mention already that it will be convenient in the proofs to restrict to the full subcategory of $[\mathbb{F}_m, \text{Set}]$ consisting of functors preserving finite connected limits. This subcategory is also known as the Schanuel topos, whose objects are nominal sets [8].

Pure λ -calculus defines such a functor Λ by $\Lambda_n = \{t \mid :; n \vdash t\}$. It satisfies the recursive equation $\Lambda_n \cong \underline{n} + \Lambda_n \times \Lambda_n + \Lambda_{n+1}$, where $- + -$ is disjoint union. In pattern unification, we consider extensions of this syntax with metavariables taking a list of distinct variables as arguments. As an example, let us add a metavariable of arity p . The extended syntax Λ' defined by $\Lambda'_n = \{t \mid M : p; n \vdash t\}$ now satisfies the recursive equation $\Lambda'_n = \underline{n} + \Lambda'_n \times \Lambda'_n + \Lambda'_{n+1} + \text{Inj}(p, n)$, where $\text{Inj}(p, n)$ is the set of injections between the cardinal sets p and n , corresponding to a choice of arguments for the metavariable. In fact, $\text{Inj}(p, n)$ is nothing but the set of morphisms between p and n in the category \mathbb{F}_m , which we denote by $\mathbb{F}_m(p, n)$.

Obviously, the functors Λ and Λ' satisfy similar recursive equations. Denoting F the endofunctor on $[\mathbb{F}_m, \text{Set}]$ mapping X to $I + X \times X + X_{-+1}$, where I is the functor mapping n to \underline{n} , the functor Λ can be characterised as the initial algebra for

F , thus satisfying the recursive equation $\Lambda \cong F(\Lambda)$. In other words, Λ is the free F -algebra on the initial functor 0 . On the other hand, Λ' is characterised as the initial algebra for $F(-) + yp$, where yp is the (representable) functor $\mathbb{F}_m(p, -) : \mathbb{F}_m \rightarrow \mathbf{Set}$, thus satisfying the recursive equation $\Lambda' \cong F(\Lambda') + yp$. In other words, Λ' is the free F -algebra on yp . Denoting T the free F -algebra monad, Λ is $T(0)$ and Λ' is $T(yp)$. Similarly, the functor $T(yp + yq)$ corresponds to extending the syntax with another metavariable of arity q .

In the view to abstracting pattern unification, these observations motivate considering functor categories $[\mathcal{A}, \mathbf{Set}]$, where \mathcal{A} is a small category where all morphisms are monomorphic (to account for the pattern condition enforcing that metavariable arguments are distinct variables), together with an endofunctor⁴ F on it. Then, the abstract definition of a syntax extended with metavariables is the free F -algebra monad T applied to a finite coproduct of representable functors.

To understand how a unification problem is stated in this general setting⁵, let us first provide the familiar metavariable context notation with a formal meaning.

Notation 1. Given a metavariable context $\Gamma = (M_1 : m_1, \dots, M_p : m_p)$, we denote the finite coproduct $\coprod_{i \in \{1, \dots, p\}} ym_i$ by $\underline{\Gamma}$, or just by Γ when the context is clear.

If $\Gamma = (M_1 : m_1, \dots, M_p : m_p)$ and Δ are metavariable contexts, a Kleisli morphism $\sigma : \Gamma \rightarrow T\Delta$ is equivalently given (by the Yoneda Lemma and the universal property of coproducts) by a λ -term $\Delta; m_i \vdash \sigma_i$ for each $i \in \{1, \dots, p\}$: this is precisely the data for a metavariable substitution $\Delta \rightarrow \Gamma$. Thus, Kleisli morphisms are nothing but metavariable substitutions. Moreover, Kleisli composition corresponds to composition of substitutions.

A unification problem can be stated as a pair of parallel Kleisli morphisms $yp \xrightarrow[t]{u} T\Gamma$ where Γ is a metavariable context, corresponding to selecting a pair of terms $\Gamma; p \vdash t, u$. A unifier is nothing but a Kleisli morphism coequalising this pair. The property required by the most general unifier means that it is the coequaliser, in the full subcategory spanned by coproducts of representable functors. The main purpose of the pattern unification algorithm consists in constructing this coequaliser, if it exists, which is the case as long as there exists a unifier, as stated in Section §III.

With this in mind, we can give categorical semantics to the unification notation in Figure 1 as follows.

Notation 2. We denote a coequaliser $A \xrightarrow[t]{u} \Gamma \xrightarrow[\sigma]{\tau} \Delta$ in a category \mathcal{B} by $\Gamma \vdash t =_{\mathcal{B}} u \Rightarrow \sigma \dashv \Delta$, sometimes even omitting \mathcal{B} .

This notation is used in the unification phase, taking \mathcal{B} to be the Kleisli category of T restricted to coproducts of

⁴In Section §II-B, we make explicit assumptions about this endofunctor for the unification algorithm to properly generalise.

⁵What follows is a generalisation of the first-order case explained in [20], [1], in the sense that we consider a free monad on a presheaf category, rather than on sets indexed by a fixed set of sorts.

representable functors, and extended with an error object \perp (as formally justified in Section §III), with the exception of the premise of the rule U-FLEXFLEX, where $\mathcal{B} = \mathcal{D}$ is the opposite category of \mathbb{F}_m . The latter corresponds to the above rule (2), whose premise precisely means that $p \xrightarrow{\bar{z}} m \xrightarrow[\bar{y}]{\bar{x}} n$ is indeed an equaliser in \mathbb{F}_m , where n is the (implicit) variable context.

Remark 3. In Notation 2, when A is a coproduct $yn_1 + \dots + yn_p$, then t and u can be thought of as lists of terms $\Gamma; n_i \vdash t_i, u_i$, hence the vector notation used in various rules (e.g., U-SPLIT). Moreover, the usage of comma as a list separator in the conclusion is formally justified by the notation $a + c \xrightarrow{f, g} b$ given morphisms $a \xrightarrow{f} b \xleftarrow{g} c$.

Note that the rule U-SPLIT is in fact valid in any category (see [20, Theorem 9]). To formally understand the rule U-NOCYCLE which we have already introduced in the case of pure λ -calculus, see (3), let us provide the pruning notation with categorical semantics.

Notation 4. We denote a pushout diagram in a category \mathcal{B} as below left by the notation as below right, sometimes even omitting \mathcal{B} .

$$\begin{array}{ccc} A & \xrightarrow{f} & \Gamma' \\ t \downarrow & & \downarrow u \\ \Gamma & \xrightarrow[\sigma]{\tau} & \Delta \end{array} \Leftrightarrow \Gamma \vdash_{\mathcal{B}} t := f \Rightarrow u; \sigma \dashv \Delta$$

Similarly to Notation 2, this is used in Figure 1, taking \mathcal{B} to be the Kleisli category of T restricted to coproducts of representable functors, and extended with an error object \perp , with the exception of the premise of the rule P-FLEX, where $\mathcal{B} = \mathcal{D}$ is the opposite category of \mathbb{F}_m . The latter corresponds to the above rule (7) whose premise precisely means that the following square is a pullback in \mathbb{F}_m .

$$\begin{array}{ccc} p & \xrightarrow{l} & n \\ r \downarrow & & \downarrow x \\ m & \xrightarrow{y} & C \end{array}$$

Remark 5. Let us add a few more comments about Notation 4. First, note that if A is a coproduct $yn_1 + \dots + yn_p$, then t and u can be thought of as lists of terms $\Gamma; n_i \vdash t_i$ and $\Gamma'; n_i \vdash u_i$. In fact, in the situations we will consider, f will be of the shape $yn_1 + \dots + yn_p \xrightarrow{f_1 + \dots + f_p} ym_1 + \dots + ym_p$. This explains our usage of $+$ as a list separator in the rule P-SPLIT.

Plan of the paper

In Section §II, we present our categorical setting. In Section §III, we state the existence of the most general unifier as a categorical property. Then we describe the construction of the most general unifier, as summarised in Figure 1, starting with the unification phase (Section §IV), the pruning phase

Unification Phase (Section §IV)

- Structural rules

$$\frac{\overline{\Gamma \vdash () = () \Rightarrow 1_\Gamma \dashv \Gamma} \quad \overline{\perp \vdash \vec{t} = \vec{u} \Rightarrow ! \dashv \perp}}{\frac{\Gamma \vdash t_1 = u_1 \Rightarrow \sigma_1 \dashv \Delta_1 \quad \Delta_1 \vdash \vec{t}_2[\sigma_1] = \vec{u}_2[\sigma_1] \Rightarrow \sigma_2 \dashv \Delta_2}{\Gamma \vdash t_1, \vec{t}_2 = u_1, \vec{u}_2 \Rightarrow \sigma_1[\sigma_2] \dashv \Delta_2}} \text{U-SPLIT}$$

- Rigid-rigid (Section §IV-A)

$$\frac{\Gamma \vdash \vec{t} = \vec{u} \Rightarrow \sigma \dashv \Delta}{\Gamma \vdash o(\vec{t}) = o(\vec{u}) \Rightarrow \sigma \dashv \Delta} \text{U-RIGRIG} \quad \frac{o \neq o'}{\Gamma \vdash o(\vec{t}) = o'(\vec{u}) \Rightarrow ! \dashv \perp}$$

- Flex-*, no cycle (Section §IV-B)

$$\frac{u|_\Gamma = u' \quad \Gamma \vdash u' :> M(x) \Rightarrow v; \sigma \dashv \Delta}{\Gamma, M : b \vdash M(x) = u \Rightarrow \sigma, M \mapsto v \dashv \Delta} \text{U-NOCYCLE} + \text{symmetric rule}$$

- Flex-Flex, same (Section §IV-C)

$$\frac{b \vdash x =_{\mathcal{D}} y \Rightarrow z \dashv c}{\Gamma, M : b \vdash M(x) = M(y) \Rightarrow M \mapsto M'(z) \dashv \Gamma, M' : c} \text{U-FLEXFLEX}$$

- Flex-Rigid, cyclic (Section §IV-D)

$$\frac{u = o(\vec{t}) \quad u|_\Gamma \neq \dots}{\Gamma, M : b \vdash M(x) = u \Rightarrow ! \dashv \perp} \text{U-CYCLIC} + \text{symmetric rule}$$

Pruning phase (Section §V)

- Structural rules

$$\frac{\overline{\Gamma \vdash () :> () \Rightarrow (); 1_\Gamma \dashv \Gamma} \quad \overline{\perp \vdash \vec{t} :> \vec{f} \Rightarrow !; ! \dashv \perp}}{\frac{\Gamma \vdash t_1 :> f_1 \Rightarrow u_1; \sigma_1 \dashv \Delta_1 \quad \Delta_1 \vdash \vec{t}_2[\sigma_1] :> \vec{f}_2 \Rightarrow \vec{u}_2; \sigma_2 \dashv \Delta_2}{\Gamma \vdash t_1, \vec{t}_2 :> f_1 + \vec{f}_2 \Rightarrow u_1[\sigma_2], \vec{u}_2; \sigma_1[\sigma_2] \dashv \Delta_2}} \text{P-SPLIT}$$

- Rigid (Section §V-A)

$$\frac{\Gamma \vdash \vec{t} :> \mathcal{L}^+ x^o \Rightarrow \vec{u}; \sigma \dashv \Delta \quad o = x \cdot o'}{\Gamma \vdash o(\vec{t}) :> N(x) \Rightarrow o'(\vec{u}); \sigma \dashv \Delta} \text{P-RIG} \quad \frac{o \neq x \cdot \dots}{\Gamma \vdash o(\vec{t}) :> N(x) \Rightarrow !; ! \dashv \perp} \text{P-FAIL}$$

- Flex (Section §V-B)

$$\frac{c \vdash_{\mathcal{D}} y :> x \Rightarrow y'; x' \dashv d}{\Gamma, M : c \vdash M(y) :> N(x) \Rightarrow M'(y'); M \mapsto M'(x') \dashv \Gamma, M' : d} \text{P-FLEX}$$

Fig. 1. Summary of the rules

(Section §V), the occur-check (Section §VI). We finally justify completeness in Section §VII.

General notations

Given $n \in \mathbb{N}$, we denote the set $\{0, \dots, n-1\}$ by \underline{n} . \mathcal{B}^{op} denotes the opposite category of \mathcal{B} . If \mathcal{B} is a category and a and b are two objects, we denote the set of morphisms between a and b by $\text{hom}_{\mathcal{B}}(a, b)$ or $\mathcal{B}(a, b)$. We denote the identity morphism at an object x by 1_x . We denote by $()$ any initial morphism and by $!$ any terminal morphism. We denote the coproduct of two objects A and B by $A+B$ and the coproduct of a family of objects $(A_i)_{i \in I}$ by $\coprod_{i \in I} A_i$, and similarly for morphisms. If $f : A \rightarrow B$ and $g : A' \rightarrow B$, we denote the induced morphism $A+A' \rightarrow B$ by f, g . Coproduct injections $A_i \rightarrow \coprod_{i \in I} A_i$ are typically denoted by in_i . Let T be a monad on a category \mathcal{B} . We denote its unit by η , and its Kleisli

category by Kl_T : the objects are the same as those of \mathcal{B} , and a Kleisli morphism from A to B is a morphism $A \rightarrow TB$ in \mathcal{B} . We denote the Kleisli composition of $f : A \rightarrow TB$ and $g : B \rightarrow TC$ by $f[g] : A \rightarrow TC$.

II. GENERAL SETTING

In our setting, syntax is specified as an endofunctor F on a category \mathcal{C} . We introduce conditions for the latter in Section §II-A and for the former in Section §II-B. Examples are presented in Section §VIII.

A. Base category

We work in a full subcategory \mathcal{C} of functors $\mathcal{A} \rightarrow \text{Set}$, namely, those preserving finite connected limits, where \mathcal{A} is a small category in which all morphisms are monomorphisms and which has finite connected limits.

Example 6. In Section §I-B, we considered $\mathcal{A} = \mathbb{F}_m$ the category of finite cardinals and injections. Note that \mathcal{C} is equivalent to the category of nominal sets [8].

Remark 7. The main property that justifies unification of two metavariables as an equaliser or a pullback in \mathcal{A} is that given any metavariable context Γ , the functor $TT : \mathcal{A} \rightarrow \text{Set}$ preserves them, i.e., $TT \in \mathcal{C}$. In fact, the argument works not only in the category of metavariable contexts and substitutions, but also in the (larger) category of objects of \mathcal{C} and Kleisli morphisms between them. However, counter-examples can be found in the total Kleisli category. Consider indeed the unification problem $M(x, y) = M(y, x)$, in the example of pure λ -calculus. We can define⁶ a functor P that does not preserve finite connected colimits such that $T(P)$ is the syntax extended with a binary commutative metavariable $M'(-, -)$. Then, the most general unifier, computed in the total Kleisli category, replaces M with P . But in the Kleisli category restricted to coproducts of representable functors, or more generally, to objects of \mathcal{C} , the coequaliser replaces M with a constant metavariable, as expected.

Remark 8. The category \mathcal{A} is intuitively the category of metavariable arities. A morphism in this category can be thought of as data to substitute a metavariable $M : a$ with another. For example, in the case of pure λ -calculus, replacing a metavariable $M : m$ with a metavariable $N : n$ amounts to a choice of distinct variables $x_1, \dots, x_n \in \{0, \dots, m-1\}$, i.e., a morphism $\text{hom}_{\mathbb{F}_m}(n, m)$. The condition that all morphisms are monomorphic can be thought of as an abstract version of the pattern restriction.

Lemma 9. \mathcal{C} is closed under limits, coproducts, and filtered colimits. Moreover, it is cocomplete.

By right continuity of the homset bifunctor, any representable functor is in \mathcal{C} and thus the embedding $\mathcal{C} \rightarrow [\mathcal{A}, \text{Set}]$ factors the Yoneda embedding $\mathcal{A}^{\text{op}} \rightarrow [\mathcal{A}, \text{Set}]$.

Notation 10. We set $\mathcal{D} = \mathcal{A}^{\text{op}}$ and denote the fully faithful Yoneda embedding as $\mathcal{D} \xrightarrow{K} \mathcal{C}$. We denote by $\mathcal{D}^+ \xrightarrow{K^+} \mathcal{C}$ the full subcategory of \mathcal{C} consisting of finite coproducts of objects of \mathcal{D} . Moreover, we adopt Notation 1 for objects of \mathcal{D}^+ , that is, a coproduct $\coprod_{i \in \{M, N, \dots\}} K a_i$ is denoted by a (metavariable) context $M : a_M, N : a_N, \dots$

Remark 11. \mathcal{D}^+ is equivalent to the category of finite families of objects of \mathcal{A} . Thinking of objects of \mathcal{A} as metavariable arities (Remark 8), \mathcal{D}^+ can be thought of as the category of metavariable contexts.

We now abstract the situation by listing a number of properties that we will use to justify the unification algorithm.

Property 12. The following properties hold.

- (i) \mathcal{D} has finite connected colimits.
- (ii) $K : \mathcal{D} \rightarrow \mathcal{C}$ preserves finite connected colimits.
- (iii) Given any morphism $f : a \rightarrow b$ in \mathcal{D} , the morphism Kf is epimorphic.

⁶Define P_n as the set of two-elements sets of $\{0, \dots, n-1\}$.

- (iv) Coproduct injections $A_i \rightarrow \coprod_j A_j$ in \mathcal{C} are monomorphisms.
- (v) For each $d \in \mathcal{D}$, the object Kd is connected, i.e., any morphism $Kd \rightarrow \coprod_i A_i$ factors through exactly one coproduct injection $A_j \rightarrow \coprod_i A_i$.

Proof. See Appendix §A. □

Remark 13. Continuing Remark 7, unification of two metavariables as pullbacks or equalisers in \mathcal{A} crucially relies on Property 12.(ii), which holds because we restrict to functors preserving finite connected limits.

B. The endofunctor for syntax

We assume given an endofunctor F on $[\mathcal{A}, \text{Set}]$ defined by

$$F(X)_a = \prod_{o \in O_a} \prod_{j \in J_{o,a}} X_{L_{o,j,a}},$$

for some functors $O : \mathcal{A} \rightarrow \text{Set}$, $J : (\int O)^{\text{op}} \rightarrow \mathbb{F}$ and $L : (\int J)^{\text{op}} \rightarrow \mathcal{A}$, where

- \mathbb{F} is the category of finite cardinals and any morphisms between them;
- $\int O$ denotes the category of elements of O whose objects are pairs of an object a of \mathcal{A} and an element o in O_a , and morphisms between (a, o) and (a', o') are morphisms $f : a \rightarrow a'$ such that $O_f(o) = o'$;
- $\int J$ denotes the category of elements of $\int O \xrightarrow{J} \mathbb{F} \hookrightarrow \text{Set}$. Objects are triples (a, o, j) , where a is an object of \mathcal{A} , $o \in O_a$, and $j \in \{0, \dots, J_{a,o} - 1\}$, and a morphism in $\int J$ between (a, o, j) and (a', o', j') is a morphism $f : a \rightarrow a'$ such that $o = O_f(o')$ and $j' = J_f(j)$.

Example 14. For pure λ -calculus where $\mathcal{A} = \mathbb{F}_m$, we have $O_n = \{a, l\} + \{v_i | 0 \leq i < n\}$, and $J_{v_i} = 0$, $J_a = 2$, $J_l = 1$, and $L_{n,o,j} = n + 1$ is if $o = l$, or n otherwise.

We moreover assume that F restricts as an endofunctor on \mathcal{C} , i.e., that it maps functors preserving finite connected limits to functors preserving finite connected limits. This has the following consequence.

Remark 15. F induces a polynomial functor on the category of sets indexed by the objects of \mathcal{A} .

Lemma 16. O preserves finite connected limits.

Proof. O is isomorphic to $F(1)$, where 1 is the constant functor mapping everything to the singleton set $\{0\}$. Since 1 trivially preserves limits, it is in \mathcal{C} and thus $F(1) \cong O$ also is. □

Lemma 17. F is finitary and generates a free monad that restricts to a monad T on \mathcal{C} . Moreover, TX is the initial algebra of $Z \mapsto X + FZ$, as an endofunctor on \mathcal{C} .

Proof. F is finitary because filtered colimits commute with finite limits [14, Theorem IX.2.1] and colimits. The free monad construction is due to [18]. □

We will be mainly interested in coequalisers in the Kleisli category restricted to objects of \mathcal{D}^+ .

Notation 18. Let $Kl_{\mathcal{D}^+}$ denote the full subcategory of Kl_T consisting of objects in \mathcal{D}^+ . Moreover, we denote by $\mathcal{L}^+ : \mathcal{D}^+ \rightarrow Kl_{\mathcal{D}^+}$ the functor which is the identity on objects and postcomposes any morphism $A \rightarrow B$ by $\eta_B : B \rightarrow TB$, and by \mathcal{L} the functor $\mathcal{D} \hookrightarrow \mathcal{D}^+ \xrightarrow{\mathcal{L}^+} Kl_{\mathcal{D}^+}$.

Property 19. The functor $\mathcal{D} \xrightarrow{\mathcal{L}} Kl_{\mathcal{D}^+}$ preserves finite connected colimits.

Notation 20. Given $f \in \text{hom}_{\mathcal{D}}(a, b)$, $u : Kb \rightarrow X$, we denote $u \circ Kf$ by $f \cdot u$.

Given $a \in \mathcal{D}$, $o : Ka \rightarrow O$, we denote $\coprod_{j \in J_{a,o}} KL_{a,o,j}$ by \bar{o} . Given $f \in \text{hom}_{\mathcal{D}}(b, a)$, we denote the induced morphism $\bar{f} \cdot \bar{o} \rightarrow \bar{o}$ by f^o .

Lemma 21. For any $X \in \mathcal{C}$, a morphism $Ka \rightarrow FX$ is equivalently given by a morphism $o \in Ka \rightarrow O$, and a morphism $f : \bar{o} \rightarrow X$.

Proof. This follows from Property 12.(v). \square

Notation 22. Given $o : Ka \rightarrow O$ and $\vec{t} : \bar{o} \rightarrow TX$, we denote the induced morphism $Ka \rightarrow FTX \hookrightarrow TX$ by $o(\vec{t})$, where the first morphism $Ka \rightarrow FTX$ is induced by Lemma 21.

Let $\Gamma = (M_1 : a_1, \dots, M_p : a_p) \in \mathcal{D}^+$ and $x \in \text{hom}_{\mathcal{D}}(a, a_i)$, we denote the Kleisli composition $Ka \xrightarrow{\mathcal{L}x} Ka_i \xrightarrow{\text{in}_i} \Gamma$ by $M_i(x) \in \text{hom}_{Kl_T}(Ka, \Gamma) = \text{hom}_{\mathcal{C}}(Ka, T\Gamma)$.

Property 23. Let $\Gamma = M_1 : a_1, \dots, M_n : a_n \in \mathcal{D}^+$. Then, any morphism $u : Ka \rightarrow T\Gamma$ is one of the two mutually exclusive following possibilities:

- $M_i(x)$ for some unique i and $x : a \rightarrow a_i$,
- $o(\vec{t})$ for some unique $o : Ka \rightarrow O$ and $\vec{t} : \bar{o} \rightarrow T\Gamma$.

We say that u is flexible (flex) in the first case and rigid in the other case.

Property 24. Let $\Gamma = M_1 : a_1, \dots, M_n : a_n \in \mathcal{D}^+$ and $\sigma : \Gamma \rightarrow T\Delta$. Then, for any $o : Ka \rightarrow O$, $\vec{t} : \bar{o} \rightarrow T\Gamma$, $u : b \rightarrow a$, $i \in \{1, \dots, n\}$, $x : a \rightarrow a_i$,

$$\begin{aligned} o(\vec{t})[\sigma] &= o(\vec{t}[\sigma]) & M_i(x)[\sigma] &= x \cdot \sigma_i \\ u \cdot (o(\vec{t})) &= (u \cdot o)(\vec{t} \circ u^o) & u \cdot M_i(x) &= M(x \circ u) \end{aligned}$$

III. MAIN RESULT

The main point of pattern unification is that a pair of parallel morphisms in $Kl_{\mathcal{D}^+}$ either has no unifier, or has a coequaliser. Working with this logical disjunction is slightly inconvenient; we rephrase it in terms of a true coequaliser by freely adding a terminal object.

Definition 25. Given a category \mathcal{B} , let \mathcal{B}^* be \mathcal{B} extended freely with a terminal object.

Notation 26. We denote by \perp the freely added terminal object in \mathcal{B}^* . Recall that $!$ denotes any terminal morphism.

Adding a terminal object results in adding a terminal cocone to all diagrams. As a consequence, we have the following lemma.

Lemma 27. Let J be a diagram in a category \mathcal{B} . The following are equivalent:

- 1) J has a colimit as long as there exists a cocone;
- 2) J has a colimit in \mathcal{B}^* .

The following result is also useful.

Lemma 28. Given a category \mathcal{B} , the canonical embedding functor $\mathcal{B} \rightarrow \mathcal{B}^*$ creates colimits.

As a consequence,

- 1) whenever the colimit in $Kl_{\mathcal{D}^+}^*$ is not \perp , it is also a colimit in $Kl_{\mathcal{D}^+}$;
- 2) existing colimits in $Kl_{\mathcal{D}^+}$ are also colimits in $Kl_{\mathcal{D}^+}^*$;
- 3) in particular, coproducts in $Kl_{\mathcal{D}^+}$ (which are computed in \mathcal{C}) are also coproducts in $Kl_{\mathcal{D}^+}^*$.

The main point of pattern unification is the following result.

Theorem 29. $Kl_{\mathcal{D}^+}^*$ has coequalisers.

In the next sections, we show how the generic unification algorithm summarised in Figure 1 provides a construction of such coequalisers.

IV. UNIFICATION PHASE

In this section, we describe the main unification phase, which computes a coequaliser in $Kl_{\mathcal{D}^+}^*$. We denote a coequaliser $\coprod_i Ka_i \xrightarrow[\vec{u}]{\vec{t}} \Gamma \xrightarrow{\sigma} \Delta$ in $Kl_{\mathcal{D}^+}^*$ by $\Gamma \vdash \vec{t} = \vec{u} \Rightarrow \sigma \dashv \Delta$, following Notation 2.

Let us start with the structural rules. When $\Gamma = \perp$, the coequaliser is the terminal cocone, i.e., $\perp \vdash \vec{t} = \vec{u} \Rightarrow ! \dashv \perp$ holds. When the coproduct is empty, the coequaliser is just Γ , i.e., $\Gamma \vdash () = () \Rightarrow 1_{\Gamma} \dashv \Gamma$ holds.

Thanks to the U-SPLIT rule which constructs coequalisers sequentially (as discussed in Section §I-C), we can now focus on the rules dealing with singleton lists, that is, with coequaliser diagrams $Ka \xrightarrow[\vec{u}]{\vec{t}} T\Gamma$. By Property 23, $t, u : Ka \rightarrow T\Gamma$ are either rigid or flexible. In the next subsections, we discuss all the different mutually exclusive situations (up to symmetry):

- both t or u are rigid (Section §IV-A),
- $t = M(\dots)$ and M does not occur in u (Section §IV-B),
- t and u are $M(\dots)$ (Section §IV-C),
- $t = M(\dots)$ and M occurs deeply in u (Section §IV-D).

A. Rigid-rigid

Here we detail unification of $o(\vec{t})$ and $o'(\vec{u})$ for some $o, o' : Ka \rightarrow O$, morphisms $\vec{t} : \bar{o} \rightarrow T\Gamma$, $\vec{u} : \bar{o}' \rightarrow T\Gamma$.

Assume given a unifier $\sigma : \Gamma \rightarrow \Delta$. By Property 24, $o(\vec{t}[\sigma]) = o'(\vec{u}[\sigma])$. By Property 23, this implies that $o = o'$, $\vec{t}[\sigma] = \vec{u}[\sigma]$. Therefore, we get the following failing rule

$$\frac{o \neq o'}{\Gamma \vdash o(\vec{t}) = o'(\vec{u}) \Rightarrow ! \dashv \perp}$$

We now assume $o = o'$. Then, $\sigma : \Gamma \rightarrow \Delta$ is a unifier if and only if it unifies \vec{t} and \vec{u} . This induces an isomorphism between

the category of unifiers for $o(\vec{t})$ and $o(\vec{u})$ and the category of unifiers for \vec{t} and \vec{u} , justifying the rule U-RIGRIG.

B. Flex-*, no cycle

Here we detail unification of $M(x)$, which is nothing but $\mathcal{L}x[in_M]$, and $u : Ka \rightarrow T(\Gamma, M : b)$, such that M does not occur in u , in the sense that $u = u'[in_\Gamma]$ for some $u' : Ka \rightarrow T\Gamma$. We exploit the following general lemma, recalling Notation 4.

Lemma 30 ([2], Exercise 2.17.1). *In any category, denoting morphism composition $g \circ f$ by $f[g]$, the following rule applies:*

$$\frac{\Gamma \vdash t :> t' \Rightarrow v; \sigma \vdash \Delta}{\Gamma + B \vdash t[in_1] = t'[in_2] \Rightarrow \sigma, v \vdash \Delta}$$

Taking $t = M(x) = \mathcal{L}x : Ka \rightarrow (M : b)$ and $t' = u'$, we thus have the rule

$$\frac{\Gamma \vdash u' :> M(x) \Rightarrow v; \sigma \vdash \Delta \quad u = u'[in_\Gamma]}{\Gamma, M : b \vdash M(x) = u \Rightarrow \sigma, M \mapsto v \vdash \Delta} \quad (8)$$

Let us make the factorisation assumption about u more effective. We can define by recursion a partial morphism from $T(\Gamma, M : b)$ to $T\Gamma$ that tries to compute u' from an input data u .

Lemma 31. *There exists $m_{\Gamma,b} : T(\Gamma, M : b) \rightarrow T\Gamma + 1$ such that a morphism $u : Ka \rightarrow T(\Gamma, M : b)$ factors as $Ka \xrightarrow{u'} T\Gamma \hookrightarrow T(\Gamma, M : b)$ if and only if $m_{\Gamma,b} \circ u = in_1 \circ u'$.*

Proof. We construct m by recursion, by equipping $T\Gamma + 1$ with an adequate F -algebra. Considering the embedding $(\Gamma, M : b) \xrightarrow{\eta+!} T\Gamma + 1$, we then get the desired morphism by universal property of $T(\Gamma, M : b)$ as a free F -algebra. The claimed property is proven by induction. \square

Therefore, we can rephrase (8) as the rule U-NOCYCLE in Figure 1, using the following notations (in the second one, we take Γ as the empty context).

Notation 32. *Given $u : Ka \rightarrow T(\Gamma, M : b)$, we denote $m_{\Gamma,b} \circ u$ by $u|_\Gamma$. Moreover, for any $u' : Ka \rightarrow T\Gamma$, we denote $in_1 \circ u' : Ka \rightarrow T\Gamma + 1$ by u' .*

Notation 33. *Let Γ and Δ be metavariable contexts and $a \in \mathcal{D}$. Any $t : Ka \rightarrow T(\Gamma + \Delta)$ induces a Kleisli morphism $(\Gamma, M : a) \rightarrow T(\Gamma + \Delta)$ which we denote by $M \mapsto t$.*

C. Flex-Flex, same metavariable

Here we detail unification of $M(x) = \mathcal{L}x[in_M]$ and $M(y) = \mathcal{L}y[in_M]$, with $x, y \in \text{hom}_{\mathcal{D}}(a, b)$. We exploit the following lemma with $u = \mathcal{L}x$ and $v = \mathcal{L}y$.

Lemma 34. *In any category, denoting morphism composition $g \circ f$ by $f[g]$, the following rule applies:*

$$\frac{B \vdash u = v \Rightarrow h \vdash C}{B + D \vdash u[in_B] = v[in_B] \Rightarrow h + 1_D \vdash C + D}$$

It follows that it is enough to compute the coequaliser of $\mathcal{L}x$ and $\mathcal{L}y$. Furthermore, by Property 12.(i) and Property 19, it can be computed as the image of the coequaliser of x and y , thus justifying the rule U-FLEXFLEX, using Notation 33.

D. Flex-rigid, cyclic

Here we handle unification of $M(x)$ for some $x \in \text{hom}_{\mathcal{D}}(a, b)$ and $u : Ka \rightarrow \Gamma, M : b$, such that u is rigid and M occurs in u , i.e., $\Gamma \rightarrow \Gamma, M : b$ does not factor u . In Section §VI, we show that in this situation, there is no unifier. Then, Lemma 31 and Notation 32 justify the rule U-CYCLIC.

V. PRUNING PHASE

The pruning phase computes a pushout in $Kl_{\mathcal{D}}^*$ of a span $\Gamma \xleftarrow{\vec{t}} \coprod_i Ka_i \xrightarrow{\coprod_i \mathcal{L}x_i} \coprod_i Kb_i$. We always implicitly assume (and enforce) that the right branch is a finite coproduct of free morphisms.

Remark 35. A pushout cocone for the above span consists in morphisms $\Gamma \xrightarrow{\sigma} T\Delta \xleftarrow{\vec{u}} \coprod_i Kb_i$ such that $\vec{t}[\sigma] = \vec{u} \circ \coprod_i Kx_i$, i.e., $t_i[\sigma] = \vec{x}_i \cdot u_i$ for each i .

When $\Gamma \xrightarrow{\sigma} T\Delta \xleftarrow{\vec{u}} \coprod_i Kb_i$ is a pushout of the above span, we use Notation 4 and denote such a situation by $\Gamma \vdash \vec{t} :> \coprod_i \mathcal{L}x_i \Rightarrow \vec{u}; \sigma \vdash \Delta$.

Let us start with the structural rules. When $\Gamma = \perp$, the pushout is the terminal cocone, i.e., $\perp \vdash \vec{t} :> \vec{f} \Rightarrow !; ! \vdash \perp$ holds. When the coproduct is empty, the pushout is just Γ , i.e., $\Gamma \vdash () :> () \Rightarrow (); 1_\Gamma \vdash \Gamma$ holds. Finally, let us note that the sequential rule P-SPLIT is in fact valid in any category.

Lemma 36. *In any category, denoting morphism composition $f \circ g$ by $g[f]$, the following rule applies.*

$$\frac{\Gamma \vdash t_1 :> f_1 \Rightarrow u_1; \sigma_1 \vdash \Delta_1 \quad \Delta_1 \vdash t_2[\sigma_1] :> f_2 \Rightarrow u_2; \sigma_2 \vdash \Delta_2}{\Gamma \vdash t_1, t_2 :> f_1 + f_2 \Rightarrow u_1[\sigma_2], u_2; \sigma_1[\sigma_2] \vdash \Delta_2} \text{P-SPLIT}$$

Thanks to the above rule, we can now focus on the case where \vec{t} is a singleton list, thus dealing with a span $T\Gamma \xleftarrow{t} Ka \xrightarrow{N(x)} T(N : b)$. By Property 23, the left morphism $Ka \rightarrow T\Gamma$ is either flexible or rigid. Each case is handled separately in the following subsections.

A. Rigid

Here, we describe the construction of a pushout of $\Gamma \xleftarrow{o(\vec{t})} Ka \xrightarrow{N(x)} N : b$ where $o : Ka \rightarrow O$ and $\vec{t} : \vec{o} \rightarrow T\Gamma$. By Remark 35, a cocone is a cospan $T\Gamma \xrightarrow{\sigma} T\Delta \xleftarrow{t'} Kb$ such that $o(\vec{t})[\sigma] = x \cdot t'$. By Property 24, this means that $o(\vec{t}[\sigma]) = x \cdot t'$. By Property 23, t' is either some $M(y)$ or $o'(\vec{u})$. But in the first case, $x \cdot t' = x \cdot M(y) = M(y \circ x)$ by Property 24, so it cannot equal $o(\vec{t}[\sigma])$, by Property 23. Therefore, $t' = o'(\vec{u})$ for some $o' : Kb \rightarrow O$ and $\vec{u} : \vec{o'} \rightarrow T\Delta$. By Property 24, $x \cdot t' = (x \cdot o')(\vec{u} \circ x^o)$. By Property 23, $o = x \cdot o'$, and $\vec{t}[\sigma] = \vec{u} \circ x^o$.

Note that there is at most one o' such that $o = x \cdot o'$, by Property 12.(iii). In this case, it follows from the above observations that a cocone is equivalently given by a cospan $T\Gamma \xrightarrow{\sigma} T\Delta \xleftarrow{\vec{u}} b^o$ such that $\vec{t}[\sigma] = \vec{u} \circ x^o$. But, by Remark 35, this is precisely the data for a pushout cocone for $\Gamma \xleftarrow{\vec{t}} a^o \xrightarrow{\mathcal{L}^+ x^o} b^o$. This actually induces an isomorphism

between the two categories of cocones, thus justifying the rules P-RIG and P-FAIL.

B. Flex

Here, we construct the pushout of $(\Gamma, M : c) \xleftarrow{M(y)} Ka \xrightarrow{N(x)} N : b$. Note that in this span, $N(x) = \mathcal{L}x$ while $M(y) = \mathcal{L}y[in_M]$. Thanks to the following lemma, it is actually enough to compute the pushout of $\mathcal{L}x$ and $\mathcal{L}y$.

Lemma 37. *In any category, denoting morphism composition by $f \circ g = g[f]$, the following rule applies*

$$\frac{X \vdash g :> f \Rightarrow u; \sigma \vdash Z}{X + Y \vdash g[in_1] :> f \Rightarrow u[in_1]; \sigma + Y \vdash Z + Y}$$

By Property 12.(i) and Property 19, the pushout of $\mathcal{L}x$ and $\mathcal{L}y$ is the image by \mathcal{L} of the pushout of x and y , thus justifying the rule P-FLEX.

VI. OCCUR-CHECK

The occur-check allows to jump from the main unification phase (Section §IV) to the pruning phase (Section §V), whenever the metavariable appearing at the top-level of the l.h.s does not occur in the r.h.s. This section is devoted to the proof that if there is a unifier of $M(\vec{x})$ and t , then either M does not occur in t , or it occurs at top-level (see Corollary 42). The argument formalises the basic intuition that $t = u[M \mapsto t]$ is impossible if M occurs deeply in u because the sizes of both hand sides can never match. To make this statement precise, we need some recursive definitions and properties of size, formally justified by exploiting the universal property of TX as the free F -algebra on X .

Definition 38. The size $|t| \in \mathbb{N}$ of a morphism $t : Ka \rightarrow T\Gamma$ is recursively defined by $|M(x)| = 0$ and $|o(\vec{t})| = 1 + |\vec{t}|$, with $|\vec{t}| = \sum_i t_i$, for any $\vec{t} : \coprod_i Ka_i \xrightarrow{\dots, t_i, \dots} T\Gamma$.

We will also need to count the occurrences of a metavariables in a term.

Definition 39. For each morphism $t : Ka \rightarrow T(\Gamma, M : b)$ we define $|t|_M$ recursively by $|M(x)|_M = 1$, $|N(x)|_M = 0$ if $N \neq M$, and $|o(\vec{t})|_M = |\vec{t}|_M$ with the sum convention as above for $|\vec{t}|_M$.

Lemma 40. *For any $t : Ka \rightarrow T(\Gamma, M : b)$, if $|t|_M = 0$, then $T\Gamma \hookrightarrow T(\Gamma, M : b)$ factors t . Moreover, for any $\Gamma = (M_1 : a_1, \dots, M_n : a_n)$, $t : Ka \rightarrow T\Gamma$, and $\sigma : \Gamma \rightarrow T\Delta$, we have $|t[\sigma]| = |t| + \sum_i |t|_{M_i} \times |\sigma_i|$.*

Corollary 41. *For any $t : Ka \rightarrow T(\Gamma, M : b)$, $\sigma : \Gamma \rightarrow T\Delta$, $x \in \text{hom}_{\mathcal{D}}(a, b)$, $u : Kb \rightarrow T\Delta$, we have $|t[\sigma, u]| \geq |t| + |u| \times |t|_M$ and $|M(x)[u]| = |u|$.*

Corollary 42. *If there is a commuting square in Kl_T*

$$\begin{array}{ccc} Ka & \xrightarrow{t} & \Gamma, M : b \\ \downarrow M(x) & & \downarrow \sigma, u \\ M : b & \xrightarrow{u} & \Delta \end{array}$$

then either $t = M(y)$ for some y , or $T\Gamma \hookrightarrow T(\Gamma, M : b)$ factors t .

Proof. Since $t[\sigma, u] = M(x)[u]$, we have $|t[\sigma, u]| = |M(x)[u]|$. Corollary 41 implies $|u| \geq |t| + |u| \times |t|_M$. Therefore, either $|t|_M = 0$ and we conclude by Lemma 40, or $|t|_M = 1$ and $|t| = 0$ and so t is $M(y)$ for some y . \square

VII. COMPLETENESS

Each inductive rule presented so far provides an elementary step for the construction of coequalisers. We need to ensure that this set of rules allows to construct a coequaliser in a finite number of steps. To make the argument more straightforward, we explicitly assume that in the splitting rules U-SPLIT and P-SPLIT in figure 1, the expressions with vector notation are not empty lists.

The following two properties are sufficient to ensure that applying rules eagerly eventually leads to a coequaliser: *progress*, i.e., there is always one rule that applies given some input data, and *termination*, i.e., there is no infinite sequence of rule applications. In this section, we sketch the proof of the latter termination property, following a standard argument. Roughly, it consists in defining the size of an input and realising that it strictly decreases in the premises. This relies on the notion of the size $|\Gamma|$ of a context Γ (as an element of \mathcal{D}^+), which can be defined as its size as a finite family of elements of \mathcal{A} (see Remark 11). We extend this definition to the case where $\Gamma = \perp$, by taking $|\perp| = 0$. We also define the size $||t||$ of a term $t : Ka \rightarrow T\Gamma$ as in Definition 38 except that we assign a size of 1 to metavariables, so that no term is of empty size.

Let us first quickly justify termination of the pruning phase. We define the size of a judgment $\Gamma \vdash f :> g \Rightarrow u; \sigma \vdash \Delta$ as $||f||$. It is straightforward to check that the sizes of the premises are strictly smaller than the size of the conclusion, for the two recursive rules P-SPLIT and P-RIG of the pruning phase, thanks to the following lemmas

Lemma 43. *For any $t : Ka \rightarrow T\Gamma$ and $\sigma : \Gamma \rightarrow T\Delta$, if σ is a renaming, i.e., $\sigma = \mathcal{L}^+ \sigma'$, for some σ' , then $||t[\sigma]|| = ||t||$.*

Lemma 44. *If there is a finite derivation tree of $\Gamma \vdash f :> g \Rightarrow u; \sigma \vdash \Delta$ and $\Delta \neq \perp$, then $|\Gamma| = |\Delta|$ and σ is a renaming.*

Now, we tackle termination for the unification phase. We define the size of a judgment $\Gamma \vdash t = u \Rightarrow \sigma \vdash \Delta$ to be the pair $(|\Gamma|, ||t||)$. The following lemma ensures that for the two recursive rules U-SPLIT and U-RIGRIG in the unification phase, the sizes of the premises are strictly smaller than the size of the conclusion, for the lexicographic order.

Lemma 45. *If there is a finite derivation tree of $\Gamma \vdash t = u \Rightarrow \sigma \vdash \Delta$, then $|\Gamma| \geq |\Delta|$, and moreover if $|\Gamma| = |\Delta|$ and $\Delta \neq \perp$, then σ is a renaming.*

VIII. APPLICATIONS

In the examples, we motivate the definition of the category \mathcal{A} based on what we expect from metavariable arities, following Remark 8.

In all our examples, O is a coproduct $\coprod_{\ell \in V} O_\ell$ for some $O_\ell : \mathcal{A} \rightarrow \text{Set}$ and $J_{a, \text{in}_\ell(o)} = \gamma_\ell$ for some finite cardinal γ_ℓ . In this case, L is equivalently given by a family of functors $H_{\ell, j} : \int O_\ell \rightarrow \text{Set}$ for each $j \in \gamma_\ell$ by $L_{a, \text{in}_\ell(o), j} = H_{\ell, j}(a, o)$. Then,

$$F(X)_a \cong \prod_{\ell \in V} \prod_{o \in O_\ell} \prod_{j \in \gamma_\ell} X_{H_{\ell, j}(a, o)}$$

A. Simply-typed second-order syntax

In this section, we present the example of simply-typed λ -calculus. Our treatment generalises to any second-order binding signature (see [6]).

Let T denote the set of simple types generated by a set of simple types. A metavariable arity $\tau_1, \dots, \tau_n \vdash \tau_f$ is given by a list of input types τ_1, \dots, τ_n and an output type τ_f . Substituting a metavariable $M : (\Gamma \vdash \tau)$ with another $M' : (\Gamma' \vdash \tau')$ requires that $\tau = \tau'$ and involves an injective renaming $\Gamma \rightarrow \Gamma'$. Thus, we consider $\mathcal{A} = \mathbb{F}_m[T] \times T$, where $\mathbb{F}_m[T]$ is the category of finite lists of elements of T and injective renamings between them.

Table I summarises the definition of the endofunctor F on $[\mathcal{A}, \text{Set}]$ specifying the syntax, where $|\Gamma|_\tau$ denotes the number (as a cardinal set) of occurrences of τ in Γ .

B. Arguments as sets

If we think of the arguments of a metavariable as specifying the available variables, then it makes sense to assemble them in a set rather than in a list. This motivates considering the category $\mathcal{A} = \mathbb{I}$ whose objects are natural numbers and a morphism $n \rightarrow p$ is a subset of $\{0, \dots, p-1\}$ of cardinal n . For instance, \mathbb{I} can be taken as subcategory of \mathbb{F}_m consisting of strictly increasing injections, or as the subcategory of the augmented simplex category consisting of injective functions. Again, we can define the endofunctor for λ -calculus as in Section §I-B. Then, a metavariable takes as argument a set of available variables, rather than a list of distinct variables. In this approach, unifying two metavariables (see the rules U-FLEXFLEX and P-FLEX) amount to computing a set intersection.

C. Quantum λ -calculus

In this section we explain how we can define pattern unification for quantum λ -calculus [16]. We denote by S the set of types, which is inductively generated as follows

$$A, B, C \in S ::= \text{qubit} | A \multimap B | !(A \multimap B) | 1 | A \otimes B | A + B | A^\ell$$

where A^ℓ is intuitively the type of finite lists of elements of type A .

We consider metavariable arities of the shape $\Delta \vdash A$, where Δ is the multiset of the argument types, and A is the output type of the metavariable. We denote by $!\Delta$ the non linear part of Δ , i.e., its sub-multiset consisting of its non-linear types, that is, types of the shape $!A$. We denote by Δ the linear part of Δ . Substituting a metavariable of arity $\Delta_1 \vdash A_1$ with a metavariable of arity $\Delta_2 \vdash A_2$ requires that $A_1 = A_2$, $\Delta_1 = \Delta_2$, and an injective renaming $!\Delta_1 \hookrightarrow !\Delta_2$. Therefore,

we choose \mathcal{A} to be the category whose objects are metavariable arities $\Delta \vdash A$ and the set of morphisms between $\Delta_1 \vdash A_1$ and $\Delta_2 \vdash A_2$ is empty if $A_1 \neq A_2$ or $\Delta_1 \neq \Delta_2$, or is the set of injective renamings between $!\Delta_1$ and $!\Delta_2$ otherwise.

The components of the endofunctor F on $[\mathcal{A}, \text{Set}]$ are specified in Table II, except for the promotion, which we discuss below. We use the following notations.

Notation 46. $\delta(P)$ denotes either a singleton set or the empty set, depending on whether the property P is true.

$|\Delta|_C$ denotes the number of occurrences of $!C$ in Δ . By convention, we take it to be 0 if $!C$ does not make sense (i.e., when C is not $A \multimap B$).

The first rule handles the term constants in $\mathcal{C} = \{\text{skip}, \text{split}^A, \text{meas}, \text{new}, U\}$, that all have typing rules of the following shape

$$\frac{c \in \mathcal{C}}{!\Delta \vdash c : A_c \multimap B_c}$$

Let us now discuss promotion for values.

$$\frac{!\Delta \vdash V : A \multimap B}{!\Delta \vdash V : !(A \multimap B)} p$$

This typing rule can be split as described in Table I, depending on what V is: a variable, a λ -abstraction, or a term constant $c \in \mathcal{C} = \{\text{skip}, \text{split}^A, \text{meas}, \text{new}, U\}$.

D. Intrinsic polymorphic syntax

We present intrinsic system F, following [12]. Let $S : \mathbb{F}_m \rightarrow \text{Set}$ mapping n to the set S_n of types for system F taking free type variables in $\{0, \dots, n-1\}$. Intuitively, a metavariable arity $n; \sigma_1, \dots, \sigma_p \vdash \tau$ specifies the number n of free type variables, the list of input types $\vec{\sigma}$, and the output type τ , all living in S_n . Substituting a metavariable $M : (n; \vec{\sigma} \vdash \tau)$ with another $M' : (n'; \vec{\sigma}' \vdash \tau')$ requires a choice $(\alpha_0, \dots, \alpha_{n-1})$ of n distinct type variables among $\{0, \dots, n'-1\}$, such that $\tau[\vec{\alpha}] = \tau'$, and an injective renaming $\vec{\sigma}'[\vec{\alpha}] \rightarrow \vec{\sigma}$. We therefore consider the category \mathcal{A} of metavariable arities where a morphism between $n; \Gamma \vdash \tau$ and $n'; \Gamma' \vdash \tau'$ is a morphism $\sigma : n \rightarrow n'$ in \mathbb{F}_m such that $\tau[\sigma] = \tau'$, and a renaming $\Gamma[\sigma] \rightarrow \Gamma'$. More formally, \mathcal{A} is the op-lax colimit of $n \mapsto \mathbb{F}_m[S_n] \times S_n$. The intrinsic syntax of system F can then be specified as in Table I.

To understand how unification of two metavariables works (see the rules U-RIGRIG and P-RIG), let us explain how finite connected limits are computed in \mathcal{A} .

Let us introduce the category \mathcal{A}' whose definition follows that of \mathcal{A} , but without the output types: objects are pairs of a natural number n and an element of S_n . Note that this is op-lax colimit of $n \mapsto \mathbb{F}_m[S_n]$, and there is an obvious projection $\mathcal{A} \rightarrow \mathcal{A}'$, which creates finite limits, as we will show.

TABLE I
EXAMPLES OF ENDOFUNCTORS FOR SYNTAX

$$F(X)_a \cong \coprod_{\ell \in V} \coprod_{o \in O_\ell} \coprod_{j \in \gamma_\ell} X_{H_{\ell,j}(a,o)}$$

Typing rule	$O_\ell(\Gamma \vdash \tau)$	$H_{\ell,j}(\Gamma \vdash \tau)$
$\frac{x : \tau \in \Gamma}{\Gamma \vdash x : \tau}$	$\{v_i i \in \Gamma _\tau\}$	-
$\frac{\Gamma \vdash t : \tau' \Rightarrow \tau \quad \Gamma \vdash u : \tau'}{\Gamma \vdash t u : \tau}$	$\{a_{\tau'} \tau' \in T\}$	$H_{-,0} = \Gamma \vdash \tau' \Rightarrow \tau$ $H_{-,1} = \Gamma \vdash \tau'$
$\frac{\Gamma, x : \tau_1 \vdash t : \tau_2}{\Gamma \vdash \lambda x. t : \tau_1 \Rightarrow \tau_2}$	$\{l_{\tau_1, \tau_2} \tau = (\tau_1 \Rightarrow \tau_2)\}$	$H_{-,0} = \Gamma, \tau_1 \vdash \tau_2$

Simply-typed λ -calculus
(Section §VIII-A)

Typing rule	$O_\ell(n; \Gamma \vdash \tau)$	$H_{\ell,j}(n; \Gamma \vdash \tau, o)$
$\frac{x : \tau \in \Gamma}{n; \Gamma \vdash x : \tau}$	$\{v_i i \in \Gamma _\tau\}$	-
$\frac{n; \Gamma \vdash t : \tau' \Rightarrow \tau \quad n; \Gamma \vdash u : \tau'}{n; \Gamma \vdash t u : \tau}$	$\{a_{\tau'} \tau' \in S_n\}$	$H_{-,0} = n; \Gamma \vdash \tau' \Rightarrow \tau$ $H_{-,1} = n; \Gamma \vdash \tau'$
$\frac{n; \Gamma, x : \tau_1 \vdash t : \tau_2}{n; \Gamma \vdash \lambda x. t : \tau_1 \Rightarrow \tau_2}$	$\{l_{\tau_1, \tau_2} \tau = (\tau_1 \Rightarrow \tau_2)\}$	$H_{-,0} = n; \Gamma, \tau_1 \vdash \tau_2$
$\frac{n; \Gamma \vdash t : \forall \tau_1 \quad \tau_2 \in S_n}{n; \Gamma \vdash t \cdot \tau_2 : \tau_1[\tau_2]}$	$\{A_{\tau_1, \tau_2} \tau = \tau_1[\tau_2]\}$	$H_{-,0} = n; \Gamma \vdash \forall \tau_1$
$\frac{n+1; wk(\Gamma) \vdash t : \tau}{n; \Gamma \vdash \Lambda t : \forall \tau}$	$\{\Lambda_{\tau'} \tau = \forall \tau'\}$	$H_{-,0} = n+1; wk(\Gamma) \vdash \tau'$

System F
(Section §VIII-D)

Typing rules for values	$O_\ell(\Delta \vdash C)$	$H_{\ell,j}(\Delta \vdash C, o)$
$!\Delta, x : !(A \multimap B) \vdash x : !(A \multimap B)$	$\delta(\Delta = !\Delta) \times \Delta _C$	-
$\frac{!\Delta, x : A \vdash M : B}{!\Delta \vdash \lambda x^A. M : !(A \multimap B)}$	$\{l_{A,B,v} \Delta \vdash C = !\Delta \vdash !(A \multimap B)\}$	$H_{-,1}(l_{A,B,v}) = \Delta, A \vdash B$
$!\Delta \vdash c : !(A_c \multimap B_c)$	$\delta(\Delta \vdash C = !\Delta \vdash !(A_c \multimap B_c))$	-

Quantum λ -calculus
(Section §VIII-C)

Lemma 47. \mathcal{A}' has finite limits, and the projection functor $\mathcal{A}' \rightarrow \mathbb{F}_m$ preserves them.

Proof. The crucial point is that \mathcal{A}' is not only op-fibred over \mathbb{F}_m by construction, it is also fibred over \mathbb{F}_m . Intuitively, if $\Gamma \in \mathbb{F}_m[S_n]$ and $f : n' \rightarrow n$ is a morphism in \mathbb{F}_m , then $f_! \Gamma \in \mathbb{F}_m[S_{n'}]$ is essentially Γ restricted to elements of S_n that are in the image of S_f . Note that $f_!$ is right adjoint to $\Gamma \mapsto \Gamma[f]$, and is thus continuous. We now apply [11, Theorem 4.2 and Proposition 4.1]: each $\mathbb{F}_m[S_n]$ has those limits. \square

Lemma 48. The projection functor $\mathcal{A} \rightarrow \mathcal{A}'$ creates finite limits.

Proof. Let $d : I \rightarrow \mathcal{A}$ be a functor. We denote d_i by $n_i; \Gamma_i \vdash \tau_i$. Let $n; \Gamma$ be the limit of $i \mapsto n_i; \Gamma_i$ in \mathcal{A}' . By the previous lemma, n is the limit of $i \mapsto n_i$. Note that $S : \mathbb{F}_m \rightarrow \text{Set}$ preserves finite connected limits. Thus, we can define $\tau \in S_n$ as corresponding to the universal function $1 \rightarrow S_n$ factorising the cone $(1 \xrightarrow{\tau_i} S_{n_i})_i$.

It is easy to check that $n; \Gamma \vdash \tau$ is the limit of d . \square

More concretely, a finite connected limit of $i \mapsto n_i; \Gamma_i \vdash \tau_i$ in \mathcal{A} is computed as follows:

- 1) compute the limit n of $(n_i)_i$ in \mathbb{F}_m , denoting $p_i : n \rightarrow n_i$ is the canonical projections;
- 2) define τ as the (only) element of S_n such that $\tau[p_i] = \tau_i$
- 3) define Γ as the limit in $\mathbb{F}_m[S_n]$ of $p_{i,!} \Gamma_i$, where $p_{i,!} : \mathbb{F}_m[S_{n_i}] \rightarrow \mathbb{F}_m[S_n]$ is the reindexing functor described in the proof of Lemma 47.

This means that to unify $M(\vec{\alpha}; \vec{\tau})$ with $M(\vec{\alpha}'; \vec{\tau}')$, with M of arity p ; $\vec{u} \vdash u'$, we first need to compute the vector of common position \vec{i} between $\vec{\alpha}$ and $\vec{\alpha}'$, i.e, the largest vector $(i_1 < \dots < i_n)$ such that $\alpha_{i_j} = \alpha'_{i_j}$. Then, we consider $\vec{\sigma}$ and $\vec{\sigma}'$ so that $\vec{\sigma}[j \mapsto i_{j+1}]$ is the sub-list of $\vec{\tau}$ that only use type variables in $\alpha_{i_1}, \dots, \alpha_{i_n}$, and similarly for $\vec{\sigma}'$ and $\vec{\tau}'$. Finally, we define (t_1, \dots, t_m) as the vector of common positions between $\vec{\sigma}$ and $\vec{\sigma}'$. The most general unifier is then $M \mapsto N(\vec{i}, \vec{t})$ for a fresh metavariable N of arity n ; $\sigma_{t_1}, \dots, \sigma_{t_m} \vdash t'$, where t' is $u'[i_{j+1} \mapsto j]$.

Typing rules (Γ, Σ linear)	$O_\ell(\Delta \vdash C)$	$H_{\ell,j}(\Delta \vdash C, o)$
$c \in \mathcal{C}$ $\frac{}{\Delta \vdash c : A_c \multimap B_c}$	$\delta(\Delta \vdash C = !\Delta \vdash A_c \multimap B_c)$	-
A linear $\frac{}{\Delta, x : A \vdash x : A}$	$\delta(\Delta = C)$	-
$\Delta, x : !(A \multimap B) \vdash x : A \multimap B$ and $\frac{\Delta, x : A \vdash M : B}{\Delta \vdash \lambda x^A. M : A \multimap B} \multimap I$	$\delta(\Delta = !\Delta) \times \Delta _{!C}$	-
$\frac{\Delta, \Gamma \vdash M : A \multimap C \quad !\Delta, \Sigma \vdash N : A}{!\Delta, \Gamma, \Sigma \vdash MN : C} \multimap E$	$\{!_{A,B} C = A \multimap B\}$	$H_{-,1} = \Delta, A \vdash B$
$\frac{!\Delta, \Gamma \vdash M : 1 \quad !\Delta, \Sigma \vdash N : C}{!\Delta, \Gamma, \Sigma \vdash M; N : C} 1_E$	$\{u_{\Gamma, \Sigma} \Delta = !\Delta, \Gamma, \Sigma\}$	$H_{-,0} = !\Delta, \Gamma \vdash 1$ $H_{-,1} = !\Delta, \Sigma \vdash C$
$\frac{!\Delta, \Gamma \vdash M : A \quad !\Delta, \Sigma \vdash N : B}{!\Delta, \Gamma, \Sigma \vdash M \otimes N : A \otimes B} \otimes I$	$\{t_{A,B,\Gamma,\Sigma} \Delta \vdash C = !\Delta, \Gamma, \Sigma \vdash A \otimes B\}$	$H_{-,0} = !\Delta, \Gamma \vdash A$ $H_{-,1} = !\Delta, \Sigma \vdash B$
$\frac{!\Delta, \Gamma \vdash M : A \otimes B \quad !\Delta, \Sigma, x : A, y : B \vdash N : C}{!\Delta, \Gamma, \Sigma \vdash \text{let } x^A \otimes y^B = M \text{ in } N : C} \otimes E$	$\{t'_{A,B,\Gamma,\Sigma} \Delta = !\Delta, \Gamma, \Sigma\}$	$H_{-,0} = !\Delta, \Gamma \vdash A \otimes B$ $H_{-,1} = !\Delta, \Sigma, A, B \vdash C$
$\frac{!\Delta, \Gamma \vdash M : A}{!\Delta, \Gamma \vdash \text{in}_\ell M : A \oplus B} \oplus I^\ell$	$\{\text{inl}_{A,B} C = A \oplus B\}$	$H_{-,0} = \Delta \vdash A$
$\frac{!\Delta, \Gamma \vdash M : B}{!\Delta, \Gamma \vdash \text{in}_r M : A \oplus B} \oplus I^r$	$\{\text{inr}_{A,B} C = A \oplus B\}$	$H_{-,0} = \Delta \vdash B$
$\frac{!\Delta, \Sigma, x : A \vdash M : C \quad !\Delta, \Sigma, y : B \vdash N : C}{!\Delta, \Gamma, \Sigma \vdash \text{match } P \text{ with } (x^A : M \mid y^B : N) : C} \oplus E$	$\{m_{A,B,\Gamma,\Sigma} \Delta = !\Delta, \Gamma, \Sigma\}$	$H_{-,0} = !\Delta, \Gamma \vdash A \oplus B$ $H_{-,1} = !\Delta, \Sigma, A \vdash C$ $H_{-,2} = !\Delta, \Sigma, B \vdash C$
$\frac{!\Delta, \Gamma \vdash M : 1 \oplus (A \otimes A^\ell)}{!\Delta, \Gamma \vdash M : A^\ell} -_\ell$	$\{\text{tail}_A C = A^\ell\}$	$H_{-,0} = \Delta \vdash 1 \oplus (A \otimes A^\ell)$
$\frac{!\Delta, \Gamma, f : !(A \multimap B) \vdash N : C \quad !\Delta, f : !(A \multimap B), x : A \vdash M : B}{!\Delta, \Gamma \vdash \text{letrec } f^{A \multimap B} x = M \text{ in } N : C} \text{rec}$	$\{\text{rec}_{A,B} A, B \in S\}$	$H_{-,0} = \Delta, f : !(A \multimap B) \vdash C$ $H_{-,1} = !\Delta, f : !(A \multimap B), A \vdash B$

TABLE II
SOME COMPONENTS OF THE ENDOFUNCTOR SPECIFYING THE QUANTUM λ -CALCULUS.

REFERENCES

- [1] Barr, M., Wells, C.: Category Theory for Computing Science. Prentice-Hall, Inc., USA (1990)
- [2] Borceux, F.: Handbook of Categorical Algebra 1: Basic Category Theory. Encyclopedia of Mathematics and its Applications, Cambridge University Press (1994). <https://doi.org/10.1017/CBO9780511525858>
- [3] D., P.G.: A note on inductive generalization. Machine Intelligence **5**, 153–163 (1970). <https://cir.nii.ac.jp/crid/1573387448853680384>
- [4] De Bruijn, N.G.: Lambda-calculus notation with nameless dummies, a tool for automatic formula manipulation, with application to the church-rosser theorem. Indagationes Mathematicae **34**, 381–392 (1972)
- [5] Dunfield, J., Krishnaswami, N.R.: Sound and complete bidirectional typechecking for higher-rank polymorphism with existentials and indexed types. Proc. ACM Program. Lang. **3**(POPL), 9:1–9:28 (2019). <https://doi.org/10.1145/3290322>, <https://doi.org/10.1145/3290322>
- [6] Fiore, M.P., Hur, C.K.: Second-order equational logic. In: Proceedings of the 19th EACSL Annual Conference on Computer Science Logic (CSL 2010) (2010)
- [7] Fiore, M., Plotkin, G., Turi, D.: Abstract syntax and variable binding. In: Proc. 14th Symposium on Logic in Computer Science IEEE (1999)
- [8] Gabbay, M.J., Pitts, A.M.: A new approach to abstract syntax involving binders. In: Proc. 14th Symposium on Logic in Computer Science IEEE (1999)
- [9] Goguen, J.A.: What is unification? - a categorical view of substitution, equation and solution. In: Resolution of Equations in Algebraic Structures, Volume 1: Algebraic Techniques. pp. 217–261. Academic (1989)
- [10] Goldfarb, W.D.: The undecidability of the second-order unification problem. Theoretical Computer Science **13**(2), 225–230 (1981). [https://doi.org/10.1016/0304-3975\(81\)90040-2](https://doi.org/10.1016/0304-3975(81)90040-2), <https://www.sciencedirect.com/science/article/pii/0304397581900402>
- [11] Gray, J.W.: Fibred and cofibred categories. In: Eilenberg, S., Harrison, D.K., MacLane, S., Röhl, H. (eds.) Proceedings of the Conference on Categorical Algebra. pp. 21–83. Springer Berlin Heidelberg, Berlin, Heidelberg (1966)
- [12] Hamana, M.: Polymorphic abstract syntax via grothendieck construction (2011)
- [13] Huet, G.P.: A unification algorithm for typed lambda-calculus. Theor. Comput. Sci. **1**(1), 27–57 (1975). [https://doi.org/10.1016/0304-3975\(75\)90011-0](https://doi.org/10.1016/0304-3975(75)90011-0), [https://doi.org/10.1016/0304-3975\(75\)90011-0](https://doi.org/10.1016/0304-3975(75)90011-0)
- [14] Mac Lane, S.: Categories for the Working Mathematician. No. 5 in Graduate Texts in Mathematics, Springer, 2nd edn. (1998)
- [15] Miller, D.: A logic programming language with lambda-abstraction, function variables, and simple unification. J. Log. Comput. **1**(4), 497–536 (1991). <https://doi.org/10.1093/logcom/1.4.497>, <https://doi.org/10.1093/logcom/1.4.497>
- [16] Pagani, M., Selinger, P., Valiron, B.: Applying quantitative semantics to higher-order quantum computing. In: Jagannathan, S., Sewell, P. (eds.) The 41st Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL ’14, San Diego, CA, USA, January 20–21, 2014. pp. 647–658. ACM (2014). <https://doi.org/10.1145/2535838.2535879>, <https://doi.org/10.1145/2535838.2535879>
- [17] Plotkin, G.: An illative theory of relations. In: Cooper, R., et al. (eds.) Situation Theory and its Applications. pp. 133–146. No. 22 in CSLI Lecture Notes, Stanford University (1990)
- [18] Reiterman, J.: A left adjoint construction related to free triples. Journal of Pure and Applied Algebra **10**, 57–71 (1977)
- [19] Robinson, J.A.: A machine-oriented logic based on the resolution principle. J. ACM **12**(1), 23–41 (jan 1965). <https://doi.org/10.1145/321250.321253>, <https://doi.org/10.1145/321250.321253>
- [20] Rydeheard, D.E., Burstall, R.M.: Computational category theory. Prentice Hall International Series in Computer Science, Prentice Hall (1988)
- [21] Urban, C., Pitts, A., Gabbay, M.: Nominal unification. In: Baaz, M., Makowsky, J.A. (eds.) Computer Science Logic. pp. 513–527. Springer Berlin Heidelberg, Berlin, Heidelberg (2003)
- [22] Vezzosi, A., Abel, A.: A categorical perspective on pattern unification. RISC-Linz p. 69 (2014)
- [23] Zhao, J., d. S. Oliveira, B.C., Schrijvers, T.: A mechanical formalization of higher-ranked polymorphic type inference. Proc. ACM Program. Lang. **3**(ICFP), 112:1–112:29 (2019). <https://doi.org/10.1145/3341716>, <https://doi.org/10.1145/3341716>

APPENDIX

PROOF OF PROPERTY 12

(i) We assume that \mathcal{A} has finite connected limits. Hence, its opposite category $\mathcal{D} = \mathcal{A}^{op}$ has finite connected colimits.

(ii) Let $y : \mathcal{A}^{op} \rightarrow [\mathcal{A}, \text{Set}]$ denote the Yoneda embedding and $J : \mathcal{C} \rightarrow [\mathcal{A}, \text{Set}]$ denote the canonical embedding, so that

$$y = J \circ K. \quad (9)$$

Now consider a finite connected limit $\lim F$ in \mathcal{A} . Then,

$$\begin{aligned} \mathcal{C}(K \lim F, X) &\cong [\mathcal{A}, \text{Set}](JK \lim F, JX) \\ &\quad (J \text{ is fully faithful}) \\ &\cong [\mathcal{A}, \text{Set}](y \lim F, JX) \quad (\text{By (9)}) \\ &\cong JX(\lim F) \quad (\text{By the Yoneda Lemma.}) \\ &\cong \lim(JX \circ F) \\ &\quad (X \text{ preserves finite connected limits}) \\ &\cong \lim([\mathcal{A}, \text{Set}](yF-, JX)) \\ &\quad (\text{By the Yoneda Lemma}) \\ &\cong \lim([\mathcal{A}, \text{Set}](JKF-, JX)) \quad (\text{By (9)}) \\ &\cong \lim \mathcal{C}(KF-, X) \quad (J \text{ is full and faithful}) \\ &\cong \mathcal{C}(\text{colim } KF, X) \\ &\quad (\text{By left continuity of the hom-set bifunctor}) \end{aligned}$$

Thus, $K \lim F \cong \text{colim } KF$.

(iii) A morphism $f : a \rightarrow b$ is epimorphic if and only if the following square is a pushout [14, Exercise III.4.4]

$$\begin{array}{ccc} a & \xrightarrow{f} & b \\ f \downarrow & & \parallel \\ b & \xlongequal{\quad} & b \end{array}$$

We conclude by (ii), because all morphisms in $\mathcal{D} = \mathcal{A}^{op}$ are epimorphic by assumption.

(iv) This follows from Lemma 9, because a morphism $f : A \rightarrow B$ is monomorphic if and only if the following square is a pullback

$$\begin{array}{ccc} A & \xlongequal{\quad} & A \\ \parallel & & \downarrow f \\ A & \xrightarrow{f} & B \end{array}$$

(v) This follows from coproducts being computed pointwise (Lemma 9), and representable functors being connected, by the Yoneda Lemma.