

Generic pattern unification

Abstract—We provide a generic second-order unification for Miller’s pattern fragment. The syntax with metavariables is parameterised by a notion of signature generalising binding signatures, covering ordered λ -calculus, or (intrinsic) polymorphic syntax such as system F. The correctness of the algorithm is stated using a categorical perspective, based on the observation that the most general unifier is an equaliser in a (non “free”) multi-sorted Lawvere theory, thus generalising the case of first-order unification.

Index Terms—Unification, Category theory, Syntax

I. INTRODUCTION

Unification consists in finding a *unifier* of two terms t, u , that is a (metavariable) substitution σ such that $t[\sigma] = u[\sigma]$. Unification algorithms try to compute a most general unifier σ , in the sense that given any other unifier δ , there exists a unique δ' such that $\delta = \sigma[\delta']$. First-order unification [18] is used in ML-style type inference systems and logic programming languages such as Prolog. For more advanced type systems, where variable binding is crucially involved, one needs second-order unification [13], which is undecidable [10]. However, Miller [15] identified a decidable fragment: in so-called *pattern unification*, metavariables are allowed to take distinct variables as arguments. In this situation, we can write an algorithm that either fails in case there is no unifier, either computes the most general unifier.

Recent results in type inference, Dunfield-Krishnaswami [5], or Jinxu et al [23], include very large proofs: the former comes with a 190 page appendix, and the latter comes with a Coq proof is many thousands of lines long -- and both of these results are for tiny kernel calculi. If we ever hope to extend this kind of result to full programming languages like Haskell or OCaml, we must raise the abstraction level of these proofs, so that they are no longer linear (with a large constant) in the size of the calculus. A close examination of these proofs shows that a large part of the problem is that the type inference algorithms make use of unification, and the correctness proofs for type inference end up essentially re-establishing the entire theory of unification for each algorithm. The reason they do this is because algorithmic typing rules essentially give a first-order functional program with no abstractions over (for example) a signature for the unification algorithm to be defined over, or any axiomatic statement of the invariants the algorithmic typing rules had to maintain.

The present work is a first step towards a general solution to this problem. Our generic unification algorithm is parameterised by an abstract notion of signature, covering simply-typed second-order syntax, ordered syntax, or (intrinsic) polymorphic syntax such as system F. We focused on

Miller pattern unification, a decidable fragment of higher-order unification where metavariables can only take distinct variables as arguments. This is already a step beyond the above-cited works [23], [5] that use plain first-order unification. Moreover, this is necessary for types with binders (e.g., fixed-point operators like $\mu a. A[a]$) as well as for rich type systems like dependent types.

As an introduction, we start by presenting pattern unification in the case of pure λ -calculus in Section §II-A. The goal of this paper is to generalise it, by parameterising the algorithm by a signature specifying a syntax. Then, in Section §IV-A, we motivate our general setting and provide categorical semantics of the algorithm, by revisiting pure λ -calculus.

Related work

First-order unification has been explained from a lattice-theoretic point of view by Plotkin [3], and later categorically analysed in [19], [9], [2, Section 9.7] as coequalisers. However, there is little work on understanding pattern unification algebraically, with the notable exception of [22], working with normalised terms of simply-typed λ -calculus. The present paper can be thought of as a generalisation of their work.

Although our notion of signature has a broader scope since we are not specifically focusing on syntax where variables can be substituted, our work is closer in spirit to the presheaf approach [7] than to the nominal sets’ [8] in that everything is explicitly scoped: terms come with their support, metavariables always appear with their scope of allowed variables.

Nominal unification [21] is an alternative to pattern unification where metavariables are not supplied with the list of allowed variables. Instead, substitution can capture variables. Nominal unification explicitly deals with α -equivalence as an external relation on the syntax, and as a consequence deals with freshness problems in addition to unification problems.

Plan of the paper

In Section §IV, we present our general categorical setting. In Section §??, we state the existence of the most general unifier as a categorical property. Then we describe the construction of the most general unifier, as summarised in Figure ??, starting with the unification phase (Section §V), the non-cyclic phase (Section §VI), the occur-check (Section §??). We finally justify completeness in Section §III.

General notations

\mathcal{B}^{op} denotes the opposite category of \mathcal{B} . If \mathcal{B} is a category and a and b are two objects, we denote the set of morphisms between a and b by $\text{hom}_{\mathcal{B}}(a, b)$ or $\mathcal{B}(a, b)$. We denote the identity morphism at an object x by 1_x . We denote by $()$ any initial morphism and by $!$ any terminal morphism. We denote

the coproduct of two objects A and B by $A + B$ and the coproduct of a family of objects $(A_i)_{i \in I}$ by $\coprod_{i \in I} A_i$, and similarly for morphisms. If $f : A \rightarrow B$ and $g : A' \rightarrow B$, we denote the induced morphism $A + A' \rightarrow B$ by f, g . Coproduct injections $A_i \rightarrow \coprod_{i \in I} A_i$ are typically denoted by in_i . Let T be a monad on a category \mathcal{B} . We denote its unit by η , and its Kleisli category by Kl_T : the objects are the same as those of \mathcal{B} , and a Kleisli morphism from A to B is a morphism $A \rightarrow TB$ in \mathcal{B} . We denote the Kleisli composition of $f : A \rightarrow TB$ and $g : B \rightarrow TC$ by $f[g] : A \rightarrow TC$.

II. SUMMARY OF THE ALGORITHM

A. An example: pure λ -calculus.

Consider the syntax of pure λ -calculus extended with metavariables satisfying the pattern restriction, encoded with De Bruijn levels, rather than De Bruijn indices [4]. More formally, the syntax is inductively generated by the following inductive rules, where Γ is a metavariable context $M_1 : m_1, \dots, M_p : m_p$ specifying a metavariable symbol M_i together with its number of arguments m_i .

$$\frac{1 \leq i \leq n}{\Gamma; n \vdash i} \quad \frac{\Gamma; n \vdash t \quad \Gamma; n \vdash u}{\Gamma; n \vdash t u} \quad \frac{\Gamma; n+1 \vdash t}{\Gamma; n \vdash \lambda t}$$

$$\frac{M : m \in \Gamma \quad 1 \leq i_1, \dots, i_m \leq n \quad i_1, \dots, i_m \text{ distinct}}{\Gamma; n \vdash M(i_1, \dots, i_m)}$$

Note that the De Bruijn level convention means that the variable bound in $\Gamma; n \vdash \lambda t$ is $n+1$.

A *metavariable substitution* $\sigma : \Gamma \rightarrow \Gamma'$ assigns to each metavariable M of arity m in Γ a term $\Gamma'; m \vdash \sigma_M$. This assignation extends (through a recursive definition) to any term $\Gamma; n \vdash t$, yielding a term $\Gamma'; n \vdash t[\sigma]$. The base case is

$$M(x_1, \dots, x_n)[\sigma] = \sigma_M\{i \mapsto x_i\}, \quad (1)$$

where $\{i \mapsto x_i\}$ is variable renaming. For example, the identity substitution $1_\Gamma : \Gamma \rightarrow \Gamma$ is defined by the term $M(1, \dots, m)$ for each metavariable declaration $M : m$ in Γ . Composition of substitutions $\sigma : \Gamma_1 \rightarrow \Gamma_2$ and $\sigma' : \Gamma_2 \rightarrow \Gamma_3$ is then defined as $(\sigma[\sigma'])_M = \sigma_M[\sigma']$.

A *unifier* of two terms $\Gamma; n \vdash t, u$ is a substitution $\sigma : \Gamma \rightarrow \Gamma'$ such that $t[\sigma] = u[\sigma]$. A *most general unifier* of t and u is a unifier $\sigma : \Gamma \rightarrow \Gamma'$ that uniquely factors any other unifier $\delta : \Gamma \rightarrow \Delta$, in the sense that there exists a unique $\delta' : \Gamma' \rightarrow \Delta$ such that $\delta = \sigma[\delta']$. We denote this situation by $\Gamma \vdash t = u \Rightarrow \sigma \dashv \Gamma'$, leaving the variable context n implicit. Intuitively, the symbol \Rightarrow separates the input and the output of the unification algorithm, which either returns a most general unifier, either fails when there is no unifier at all (for example, when unifying $t_1 t_2$ with λu). To handle the latter case, we add¹ a formal error metavariable context \perp in which the only term (in any variable context) is a formal error term $!$, inducing a unique substitution $! : \Gamma \rightarrow \perp$, satisfying $t[!] = !$ for any term t . For example, we have $\Gamma \vdash t_1 t_2 = \lambda u \Rightarrow ! \dashv \perp$.

We generalise the notation (and thus the input of the unification algorithm) to lists of terms $\vec{t} = (t_1, \dots, t_n)$ and

$\vec{u} = (u_1, \dots, u_n)$ such that $\Gamma; n_i \vdash t_i, u_i$. Then, $\Gamma \vdash \vec{t} = \vec{u} \Rightarrow \sigma \dashv \Gamma'$ means that σ unifies each pair (t_i, u_i) and is the most general one, in the sense that it uniquely factors any other substitution that unifies each pair (t_i, u_i) . As a consequence, we get the *congruence* rule for application.

$$\frac{\Gamma \vdash t_1, t_2 = u_1, u_2 \Rightarrow \sigma \dashv \Delta}{\Gamma \vdash t_1 t_2 = u_1 u_2 \Rightarrow \sigma \dashv \Delta}$$

Unifying a list of term pairs $t_1, \vec{t}_2 = u_1, \vec{u}_2$ can be performed sequentially by first computing the most general unifier σ_1 of (t_1, u_1) , then applying the substitution to (\vec{t}_2, \vec{u}_2) , and finally computing the most general unifier of the resulting list of term pairs: this is precisely the rule UA-SPLIT in Figure 1.

Thanks to this rule, we can focus on unification of a single term pair. The idea here is to recursively inspect the structure of the given terms, until reaching a metavariable application $M(x_1, \dots, x_n)$ at top level on either hand side of $\Gamma, M : n \vdash t, u$. Assume by symmetry $t = M(x_1, \dots, x_n)$, then three mutually exclusive situations must be considered:

- 1) M occurs deeply in u ;
- 2) M occurs in u at top level, i.e., $u = M(y_1, \dots, y_n)$;
- 3) M does not occur in u .

In the first case, there is no unifier because the size of both hand sides can never match after substitution. This justifies the rule UA-CYCLIC, where $M \in u$ means that M occurs in u .

In the second case, we are unifying $M(\vec{x})$ with $M(\vec{y})$. The most general unifier σ coincides with the identity substitution except for $\sigma_M = M'(\vec{z})$, where $\vec{z} = (z_1, \dots, z_p)$ is the family of common positions i such that $x_i = y_i$, and M' is fresh. We denote² such a situation by $n \vdash \vec{x} = \vec{y} \Rightarrow \vec{z} \dashv p$. We therefore get the rule UA-FLEX.

Example 1. Let x, y, z be three distinct variables. The most general unifier of $M(x, y)$ and $M(z, y)$ is $M \mapsto M(2)$; the most general unifier of $M'(x, y)$ and $M(z, x)$ is $M \mapsto M'$.

Finally, the last case consists in unifying $M(\vec{x})$ with some u such that M does not occur in u , i.e., u restricts to the smaller metavariable context Γ . We denote such a situation by $u|_\Gamma = \underline{u}'$, where u' is essentially u but considered in the smaller metavariable context Γ . The algorithm then enters a *non-cyclic phase*, which specifically addresses such non-cyclic unification problems. Let us introduce a specific notation: $\Gamma \vdash u' : \triangleright M(\vec{x}) \Rightarrow w; \sigma \dashv \Delta$ means that u' is a term in a metavariable context Γ , while M is a fresh metavariable with respect to Γ and $\vec{x} = (x_1, \dots, x_m)$ are distinct variables in the (implicit) variable context of u' . The output is the most general unifier of u' and $M(\vec{x})$, both considered in the extended metavariable context $\Gamma, M : m$. This substitution from $\Gamma, M : m$ to Δ is explicitly defined as the extension of a substitution $\sigma : \Gamma \rightarrow \Delta$ with a term $\Delta; m \vdash w$ for substituting M .

²The similarity with the above introduced notation is no coincidence: as we will see, both are coequalisers.

¹We will interpret \perp as a freely added terminal object in Section §IV-A.

Judgments

$$\begin{aligned}
\Gamma \vdash \vec{t} = \vec{u} \Rightarrow \sigma \dashv \Delta &\iff \sigma : \Gamma \rightarrow \Delta \text{ is the most general unifier of } \vec{t} \text{ and } \vec{u} \\
\Gamma \vdash \vec{u} :> \overrightarrow{M(\vec{x})} \Rightarrow \vec{w}; \sigma \dashv \Delta &\iff \sigma : \Gamma \rightarrow \Delta \text{ extended with } M_i \mapsto w_i \text{ is the most general unifier of } \Gamma \vdash \vec{u} \text{ and } \overrightarrow{M(\vec{x})} \\
m \vdash \vec{x} = \vec{y} \Rightarrow \vec{z} \dashv p &\iff (z_1, \dots, z_p) \text{ are the common positions of } (x_1, \dots, x_m) \text{ and } (y_1, \dots, y_m) \\
n \vdash \vec{x} :> \vec{y} \Rightarrow \vec{l}; \vec{r} \dashv p &\iff (l_1, \dots, l_p) \text{ and } (r_1, \dots, r_p) \text{ are the common value positions of } (x_1, \dots, x_n) \text{ and } \vec{y}
\end{aligned}$$

Unification Phase

- Structural rules

$$\frac{\overline{\Gamma \vdash () = () \Rightarrow 1_\Gamma \dashv \Gamma} \quad \overline{\perp \vdash \vec{t} = \vec{u} \Rightarrow ! \dashv \perp}}{\Gamma \vdash t_1 = u_1 \Rightarrow \sigma_1 \dashv \Delta_1 \quad \Delta_1 \vdash t_2[\sigma_1] = u_2[\sigma_1] \Rightarrow \sigma_2 \dashv \Delta_2} \text{U}\Lambda\text{-SPLIT}$$

$$\frac{}{\Gamma \vdash t_1, t_2 = u_1, u_2 \Rightarrow \sigma_1[\sigma_2] \dashv \Delta_2}$$

- Rigid-rigid (o, o' are applications, λ -abstractions, or variables)

$$\frac{\Gamma \vdash \vec{t} = \vec{u} \Rightarrow \sigma \dashv \Delta}{\Gamma \vdash o(\vec{t}) = o(\vec{u}) \Rightarrow \sigma \dashv \Delta} \text{U}\Lambda\text{-RIGRIG} \quad \frac{o \neq o'}{\Gamma \vdash o(\vec{t}) = o'(\vec{u}) \Rightarrow ! \dashv \perp}$$

- Flex-*, no cycle

$$\frac{M \notin u \quad \Gamma \vdash u :> M(\vec{x}) \Rightarrow w; \sigma \dashv \Delta}{\Gamma, M : m \vdash M(\vec{x}) = u \Rightarrow \sigma, M \mapsto w \dashv \Delta} \text{U}\Lambda\text{-NOCYCLE} + \text{symmetric rule}$$

- Flex-Flex, same

$$\frac{m \vdash \vec{x} = \vec{y} \Rightarrow \vec{z} \dashv p}{\Gamma, M : m \vdash M(\vec{x}) = M(\vec{y}) \Rightarrow M \mapsto M'(\vec{z}) \dashv \Gamma, M' : p} \text{U}\Lambda\text{-FLEX}$$

- Flex-Rigid, cyclic

$$\frac{M \in u \quad u \neq M(\dots)}{\Gamma, M : m \vdash M(\vec{x}) = u \Rightarrow ! \dashv \perp} \text{U}\Lambda\text{-CYCLIC} + \text{symmetric rule}$$

Non-cyclic phase

- Structural rules

$$\frac{\overline{\Gamma \vdash () :> () \Rightarrow (); 1_\Gamma \dashv \Gamma} \quad \overline{\perp \vdash \vec{t} :> \overrightarrow{M(\vec{x})} \Rightarrow !; ! \dashv \perp}}{\Gamma \vdash t_1 :> M(\vec{x}_1) \Rightarrow u_1; \sigma_1 \dashv \Delta_1 \quad \Delta_1 \vdash \vec{t}_2[\sigma_1] :> \overrightarrow{M(\vec{x}_2)} \Rightarrow \vec{u}_2; \sigma_2 \dashv \Delta_2} \text{P}\Lambda\text{-SPLIT}$$

$$\frac{}{\Gamma \vdash t_1, \vec{t}_2 :> M_1(\vec{x}_1), \overrightarrow{M(\vec{x}_2)} \Rightarrow u_1[\sigma_2], \vec{u}_2; \sigma_1[\sigma_2] \dashv \Delta_2}$$

- Rigid

$$\frac{\Gamma \vdash t :> M'(\vec{x}, \overbrace{n+1}^{\text{bound variable}}) \Rightarrow w; \sigma \dashv \Delta}{\Gamma \vdash \lambda t :> M(\vec{x}) \Rightarrow \lambda w; \sigma \dashv \Delta} \text{P}\Lambda\text{-LAM} \quad \frac{\Gamma \vdash t, u :> M_1(\vec{x}), M_2(\vec{x}) \Rightarrow w_1, w_2; \sigma \dashv \Delta}{\Gamma \vdash t u :> M(\vec{x}) \Rightarrow w_1 w_2; \sigma \dashv \Delta} \text{P}\Lambda\text{-APP}$$

$$\frac{y = x_i}{\Gamma \vdash y :> M(\vec{x}) \Rightarrow i; 1_\Gamma \dashv \Gamma} \text{P}\Lambda\text{-VAROK} \quad \frac{y \notin \vec{x}}{\Gamma \vdash y :> M(\vec{x}) \Rightarrow !; ! \dashv \perp} \text{P}\Lambda\text{-VARFAIL}$$

- Flex

$$\frac{n \vdash \vec{x} :> \vec{y} \Rightarrow \vec{l}; \vec{r} \dashv p}{\Gamma, N : n \vdash N(\vec{x}) :> M(\vec{y}) \Rightarrow P(\vec{l}); N \mapsto P(\vec{r}) \dashv \Gamma, P : p} \text{P}\Lambda\text{-FLEX}$$

Fig. 1. Unification for pure λ -calculus (Section §II-A)

Remark 2. The symbol $:>$ evokes the *pruning* involved in this phase. Indeed, one intuition behind the non-cyclic unification of $M(\vec{x})$ and u consists in taking $u[x_i \mapsto i]$ as a definition for M . This only makes sense if the free variables of u are among \vec{x} : if u is a variable that does not occur in \vec{x} , then obviously there is no unifier. However, it is possible to remove the *outbound* variables in u if they only occur in metavariable arguments, by restricting the arities of those metavariables. We accordingly call $\sigma : \Gamma \rightarrow \Delta$ the pruning substitution. As an example, if u is a metavariable application $N(\vec{x}, \vec{y})$, then although the free variables are not all included in \vec{x} , there is still a most general unifier, and the corresponding pruning substitution essentially replaces N with M , discarding the outbound variables \vec{y} .

The previous discussion justifies the rule UA-NOCYCLE.

The non-cyclic phase proceeds recursively by introducing fresh metavariables for each argument of the top-level operation. Let us look for example at the congruence rule PA-LAM for λ -abstraction. Note that a fresh variable M' is introduced for the body of the λ -abstraction which is supplied the bound variable $n+1$ as an argument, as it should not be pruned. Keeping in mind the intuition that $M = \lambda M'$, if M' is to be substituted with w , then M should be substituted with λw , as seen in the conclusion.

As before, the rule PA-APP for application motivates generalising the non-cyclic phase to handle lists. More formally given a list $\vec{u} = (u_1, \dots, u_p)$ of terms in context $\Gamma; n_i \vdash u_i$, and lists of pruning patterns $(\vec{x}_1, \dots, \vec{x}_p)$ where each \vec{x}_i is a choice of distinct variables in n_i , the judgement $\Gamma \vdash \vec{u} :> M_1(\vec{x}_1), \dots, M_p(\vec{x}_p) \Rightarrow \vec{w}; \sigma \dashv \Delta$ means that the substitution $\sigma : \Gamma \rightarrow \Delta$ extended with $M_i \mapsto w_i$ is the most general unifier of \vec{u} and $M_1(\vec{x}_1), \dots, M_p(\vec{x}_p)$ in the extended metavariable context $\Gamma, M_1 : n_1, \dots, M_p : n_p$.

The sequential rule UA-SPLIT can be adapted to the non-cyclic phase as in the rule PA-SPLIT in Figure 1. Thanks to this sequential rule, we can focus on pruning a single term. The variable case is straightforward (rules PA-VAROK and PA-VARFAIL).

The remaining case consists in unifying $N(\vec{x})$ and $M(\vec{y})$. Consider the family z_1, \dots, z_p of common values in x_1, \dots, x_n and y_1, \dots, y_m , so that $z_i = x_{l_i} = y_{r_i}$ for some lists (l_1, \dots, l_p) and (r_1, \dots, r_p) of distinct elements of $\{1, \dots, n\}$ and $\{1, \dots, m\}$ respectively. We denote³ such a situation by $n \vdash \vec{x} :> \vec{y} \Rightarrow \vec{l}; \vec{r} \dashv p$. Then, the most general unifier replaces N with $P(\vec{r})$ for some fresh metavariable P of arity p , while the metavariable M is replaced with $P(\vec{l})$, as seen in the rule PA-FLEX.

Example 3. Let x, y, z be three distinct variables. The most general unifier of $M(x, y)$ and $N(z, x)$ is $M \mapsto N'(1), N \mapsto N'(2)$. The most general unifier of $M(x, y)$ and $N(z)$ is $M \mapsto N', N \mapsto N'$.

³Again, the similarity with the notation for non-cyclic unification is no coincidence: both are pullbacks.

This ends our description of the unification algorithm, in the specific case of pure λ -calculus. The purpose of this work is to present a generalisation, by parameterising the algorithm by a signature specifying a syntax.

B. Generalisation

Our algorithm is parameterised by a notion of signature that generalises binding signatures [16]. To recall, a binding signature (O, κ, α) specifies a set of operation symbols O together with an arity $\alpha_o = (\bar{o}_1, \dots, \bar{o}_{\kappa_o})$ which is a list of natural numbers specifying how many arguments are bound in each argument. For example, pure λ -calculus is specified by $O = \{\text{lam}, \text{app}\}$ with $\alpha_{\text{app}} = (0, 0)$ and $\alpha_{\text{lam}} = (1)$.

Our algorithm is parameterised by a *generalised binding signature*, or GB-signature, a notion we will formally introduce in Definition 19. A GB-signature consists in a tuple $(\mathcal{A}, O, \kappa, \alpha)$ consisting of

- a small category \mathcal{A} whose objects are called *arities* or *variable contexts*, and morphisms are called *renamings*;
- for each arity a , a set of operation symbols O_a ;
- for each operation symbol $o \in O_a$, a list of arities $\alpha_o = (\bar{o}_1, \dots, \bar{o}_{\kappa_o})$

such that O and α are functorial in a suitable sense (see Remark 7 below for more details). We will call an element of $\{1, \dots, \kappa_o\}$ an *o-index*.

Intuitively, O_a is the set of operation symbols available in the variable context a .

Definition 4. The syntax specified by a GB-signature (\mathcal{A}, O, α) is inductively defined by the two following rules.

$$\frac{o \in O_a \quad \Gamma; \bar{o}_1 \vdash t_1 \quad \dots \quad \Gamma; \bar{o}_{\kappa_o} \vdash t_{\kappa_o}}{\Gamma; a \vdash o(t_1, \dots, t_{\kappa_o})} \text{RIG}$$

$$\frac{M : m \in \Gamma \quad x \in \text{hom}_{\mathcal{A}}(m, a)}{\Gamma; a \vdash M(x)} \text{FLEX}$$

where a context $\Gamma; a$ consists of a variable context a and a metavariable context Γ , that is, an arity function from a finite set of metavariable symbols to the underlying set of objects of \mathcal{A} . We call a term *rigid* if it is of the shape $o(\dots)$, *flexible* if it is $M(\dots)$.

Remark 5. The syntax in the empty metavariable context does not depend on the morphisms in \mathcal{A} . In fact, by restricting the morphisms in the category of arities to identity morphisms, any GB-signature induces an indexed container [1] generating the same syntax without metavariables.

Example 6. Binding signatures can be compiled into GB-signatures. More specifically, a syntax specified by a binding signature (O, κ, α) is also generated by the GB-signature $(\mathbb{F}_m, O', \kappa', \alpha')$, where

- \mathbb{F}_m is the category of finite cardinals and injections between them;
- $O'_n = \{v_1, \dots, v_n\} \sqcup \{o_n \mid o \in O\}$;
- $\kappa'_o = \kappa_o$;
- $\alpha'_{v_i} = ()$ and $\alpha'_{o_n} = (n + \bar{o}_1, \dots, n + \bar{o}_{\kappa_o})$ for any $o \in O$, $i, n \in \mathbb{N}$.

Note that variables v_i are explicitly specified as nullary operations and thus do not require a separate formation rule, contrary to what happens with binding signatures. Moreover, the choice of renamings (i.e., morphisms in the category of arities) is motivated by the FLEX rule. Indeed, if M has arity $m \in \mathbb{N}$, then a choice of arguments in the variable context $a \in \mathbb{N}$ consists of a list of distinct variables in the variable context a , or equivalently, an injection between the cardinal sets m and a , that is, a morphism in \mathbb{F}_m between m and a .

GB-signatures also capture multi-sorted binding signatures such as simply-typed λ -calculus, or polymorphic syntaxes such as System F (see Section §VII).

Remark 7. In the notion of GB-signature, functoriality ensures that the generated syntax supports renaming: given a morphism $f : a \rightarrow b$ in \mathcal{A} and a term $\Gamma; a \vdash t$, we can recursively define a term $\Gamma; b \vdash t\{f\}$. The case of metavariables is simple: $M(x)\{f\} = M(f \circ x)$. For an operation $o(t_1, \dots, t_{\kappa_o})$, functoriality provides the following components:

- an operation symbol $o\{f\} \in O_b$;
- a list of o -indices $(\epsilon(o, f)_1, \dots, \epsilon(o, f)_{\kappa_{o\{f\}}})$;
- a morphism $f_i^o : \bar{o}_{\epsilon(o, f)_i} \rightarrow \overline{o\{f\}}_i$ for each $o\{f\}$ -index i .

Then, $o(t_1, \dots, t_{\kappa_o})\{f\}$ is defined as $o\{f\}(t_{\epsilon(o, f)_1}\{f_1^o\}, \dots)$.

Remark 8. In our examples, $\kappa_{o\{f\}}$ is always equal to κ_o and $(\epsilon(o, f)_1, \dots, \epsilon(o, f)_{\kappa_{o\{f\}}})$ is $(1, \dots, \kappa_o)$, so that f_i^o is a morphism between \bar{o}_i and $\overline{o\{f\}}_i$ and thus $o(t_1, \dots, t_{\kappa_o})\{f\} = o\{f\}(t_1\{f_1^o\}, \dots, t_{\kappa_o}\{f_{\kappa_o}^o\})$.

Unification can be formulated by following the same route as the pure λ -calculus. A metavariable substitution σ from $\Gamma = (M_1 : m_1, \dots, M_p : m_p)$ to Δ is a list of terms $(\sigma_1, \dots, \sigma_p)$ such that $\Delta; m_i \vdash \sigma_i$. Given a term $\Gamma; a \vdash t$, we can recursively define the substituted case $\Delta; a \vdash t[\sigma]$ by $o(\vec{t})[\sigma] = o(\vec{t}[\sigma])$ and $M_i(x)[\sigma] = \sigma_i\{x\}$.

Figure 2 summarises our generic algorithm, parameterised by a GB-signature, and a few more parameters: a solver for the equation

$$o = o'\{x\}, \quad (2)$$

where o' is the unknown (see the rules P-RIG and P-FAIL, detailed below), and a construction of equalisers and pullbacks in \mathcal{A} , highlighted in blue. These are used to compute the most general unifier of two metavariable applications in the rules U-FLEX and P-FLEX. For example, for pure λ -calculus, they correspond to the rules U- λ -FLEX and P- λ -FLEX: indeed, the vector of common positions of \vec{x} and \vec{y} can be characterised as their equaliser in \mathbb{F}_m , when thinking of lists as functions from a finite domain, while the vectors of common value positions can be characterised as a pullback.

The main differences with the example of pure λ -calculus presented above in Figure 1 is that the vector notation is dropped for arguments of metavariables, since they are now morphisms in a category. Moreover, the case for operations and variables are merged in the non-cyclic phase with the rules P-RIG and P-FAIL that both are handling non-cyclic unification of $M(x)$ with $o(\vec{t})$. If $o = o'\{x\}$ for some o' , then

P-RIG applies; otherwise, P-FAIL applies. Let us explain how they specialise for a syntax specified by a binding signature as in Example 6. In this case, x is a list of distinct variables (x_1, \dots, x_m) . If o is a variable, then the side condition $o = o'\{x\}$ that means that o is $x_{o'}$ for some o' . Thus, those rules specialises to the rules P- λ -VAROK and P- λ -VARFAIL. On the other hand, if o is actually o_n , i.e., an operation symbol in variable context n , then $o_n = o_m\{x\}$ and thus the rule P-RIG always applies with $o' = o_m$. Moreover, x_i^o is in fact the morphism from $m + \alpha_i$ to $n + \alpha_i$ defined as $x + \alpha_i$, corresponding to the list $(\vec{x}, n + 1, \dots, n + \alpha_i)$. Thus, the rule P-RIG unfolds as

$$\frac{\Gamma \vdash \vec{t} :> \dots, M_i(\overbrace{\vec{x}, n + 1, \dots, n + 1 + \alpha_i}^{\text{bound variables}}), \dots \Rightarrow \vec{w}; \sigma \vdash \Delta}{\Gamma \vdash o(\vec{t}) :> M(\vec{x}) \Rightarrow o(\vec{w}); \sigma \vdash \Delta}.$$

The correctness of the algorithm relies on additional assumptions on the GB-signature that we will introduce later. In particular all morphisms in \mathcal{A} must be monomorphic: this ensures that Equation (2) has at most one solution.

III. PROGRESS AND TERMINATION

Each inductive rule in Figure 2 provides an elementary step for the construction of the most general unifier. To ensure that this set of rules describes a terminating algorithm, we essentially need two properties: *progress*, i.e., there is always one rule that applies given some input data, and *termination*, i.e., there is no infinite sequence of rule applications. The former is straightforward to check, but we may an explicit argument in. In this section, we sketch the proof of the latter termination property, following a standard argument. Roughly, it consists in defining the size of an input and realising that it strictly decreases in the premises. This relies on the notion of the size $|\Gamma|$ of a metavariable context Γ , as the number of its declared metavariables. We extend this definition to the case where $\Gamma = \perp$, by taking $|\perp| = 0$. We also recursively define the size $\|t\|$ of a term t by $\|M(x)\| = 1$ and $\|o(\vec{t})\| = 1 + \|\vec{t}\|$, with $\|\vec{t}\| = \sum_i \|t_i\|$. Note that no term is of empty size.

Let us first quickly justify termination of the non-cyclic phase. We define the size of a judgment $\Gamma \vdash \vec{t} :> \overrightarrow{M(x)} \Rightarrow \vec{w}; \sigma \vdash \Delta$ as $\|\vec{t}\|$. It is straightforward to check that the sizes of the premises are strictly smaller than the size of the conclusion, for the two recursive rules P-SPLIT and P-RIG of the non-cyclic phase, thanks to the following lemmas

Lemma 9. *For any term $\Gamma; a \vdash t$ and substitution $\sigma : \Gamma \rightarrow \Delta$, if σ is a metavariable renaming, i.e., σ_M is a metavariable application for any $M : m \in \Gamma$, then $\|t[\sigma]\| = \|t\|$.*

Lemma 10. *If there is a finite derivation tree of $\Gamma \vdash \vec{t} :> \overrightarrow{M(x)} \Rightarrow \vec{w}; \sigma \vdash \Delta$ and $\Delta \neq \perp$, then $|\Gamma| = |\Delta|$ and σ is a metavariable renaming.*

Now, we tackle termination for the unification phase. We define the size of a judgment $\Gamma \vdash t = u \Rightarrow \sigma \vdash \Delta$ to be the pair $(|\Gamma|, \|t\|)$. The following lemma ensures that for the two recursive rules U-SPLIT and U-RIGRIG in the unification

Judgments

$$\begin{array}{ll}
\Gamma \vdash \vec{t} = \vec{u} \Rightarrow \sigma \dashv \Delta & \iff \sigma : \Gamma \rightarrow \Delta \text{ is the most general unifier of } \vec{t} \text{ and } \vec{u} \\
\Gamma \vdash \vec{u} :> \overrightarrow{M(x)} \Rightarrow \vec{w}; \sigma \dashv \Delta & \iff \sigma : \Gamma \rightarrow \Delta \text{ extended with } M_i \mapsto w_i \text{ is the most general unifier of } \Gamma \vdash \vec{u} \text{ and } \overrightarrow{M(x)} \\
m \vdash x = y \Rightarrow z \dashv p & \iff p \xrightarrow{z} m \xrightarrow[x]{y} \dots \text{ is an equaliser in } \mathcal{A} \\
n \vdash x :> y \Rightarrow l; r \dashv p & \iff \begin{array}{ccc} p & \xrightarrow{l} & n \\ r \downarrow & & \downarrow x \\ \dots & \xrightarrow{y} & \dots \end{array} \text{ is a pullback in } \mathcal{A}
\end{array}$$

Unification Phase

- Structural rules

$$\frac{\overline{\Gamma \vdash () = () \Rightarrow 1_\Gamma \dashv \Gamma} \quad \overline{\perp \vdash \vec{t} = \vec{u} \Rightarrow ! \dashv \perp}}{\Gamma \vdash t_1 = u_1 \Rightarrow \sigma_1 \dashv \Delta_1 \quad \Delta_1 \vdash \vec{t}_2[\sigma_1] = \vec{u}_2[\sigma_1] \Rightarrow \sigma_2 \dashv \Delta_2} \text{U-SPLIT} \\
\Gamma \vdash t_1, \vec{t}_2 = u_1, \vec{u}_2 \Rightarrow \sigma_1[\sigma_2] \dashv \Delta_2$$

- Rigid-rigid

$$\frac{\Gamma \vdash \vec{t} = \vec{u} \Rightarrow \sigma \dashv \Delta}{\Gamma \vdash o(\vec{t}) = o(\vec{u}) \Rightarrow \sigma \dashv \Delta} \text{U-RIGRIG} \quad \frac{o \neq o'}{\Gamma \vdash o(\vec{t}) = o'(\vec{u}) \Rightarrow ! \dashv \perp} \text{U-CLASH}$$

- Flex-*, no cycle

$$\frac{M \notin u \quad \Gamma \vdash u :> M(x) \Rightarrow w; \sigma \dashv \Delta}{\Gamma, M : m \vdash M(x) = u \Rightarrow \sigma, M \mapsto w \dashv \Delta} \text{U-NOCYCLE} + \text{symmetric rule}$$

- Flex-Flex, same

$$\frac{m \vdash x = y \Rightarrow z \dashv p}{\Gamma, M : m \vdash M(x) = M(y) \Rightarrow M \mapsto M'(z) \dashv \Gamma, M' : p} \text{U-FLEX}$$

- Flex-Rigid, cyclic

$$\frac{M \in u \quad u \neq M(\dots)}{\Gamma, M : m \vdash M(x) = u \Rightarrow ! \dashv \perp} \text{U-CYCLIC} + \text{symmetric rule}$$

Non-cyclic phase

- Structural rules

$$\frac{\overline{\Gamma \vdash () :> () \Rightarrow (); 1_\Gamma \dashv \Gamma} \quad \overline{\perp \vdash \vec{t} :> \overrightarrow{M(x)} \Rightarrow !; ! \dashv \perp}}{\Gamma \vdash t_1 :> M_1(x_1) \Rightarrow u_1; \sigma_1 \dashv \Delta_1 \quad \Delta_1 \vdash \vec{t}_2[\sigma_1] :> \overrightarrow{M_2(x_2)} \Rightarrow \vec{u}_2; \sigma_2 \dashv \Delta_2} \text{P-SPLIT} \\
\Gamma \vdash t_1, \vec{t}_2 :> M_1(x_1), \overrightarrow{M_2(x_2)} \Rightarrow u_1[\sigma_2], \vec{u}_2; \sigma_1[\sigma_2] \dashv \Delta_2$$

- Rigid

$$\frac{\Gamma \vdash \vec{t} :> M_1(x'_1), \dots, M_1(x'_{\kappa_o}) \Rightarrow \vec{u}; \sigma \dashv \Delta \quad o = o'\{x\}}{\Gamma \vdash o(\vec{t}) :> M(x) \Rightarrow o'(\vec{u}); \sigma \dashv \Delta} \text{P-RIG} \quad \frac{o \neq \dots \{x\}}{\Gamma \vdash o(\vec{t}) :> M(x) \Rightarrow !; ! \dashv \perp} \text{P-FAIL}$$

- Flex

$$\frac{n \vdash x :> y \Rightarrow l; r \dashv p}{\Gamma, N : n \vdash N(x) :> M(y) \Rightarrow P(l); N \mapsto P(r) \dashv \Gamma, P : p} \text{P-FLEX}$$

Fig. 2. Generic pattern unification algorithm (Section §II-B)

phase, the sizes of the premises are strictly smaller than the size of the conclusion, for the lexicographic order.

Lemma 11. *If there is a finite derivation tree of $\Gamma \vdash \vec{t} = \vec{u} \Rightarrow \sigma \dashv \Delta$, then $|\Gamma| \geq |\Delta|$, and moreover if $|\Gamma| = |\Delta|$ and $\Delta \neq \perp$, then σ is a renaming.*

IV. CATEGORICAL SEMANTICS

It remains to be proven that each rule is sound, e.g., for the rule U-SPLIT, if the output of the premises are most general unifiers, then so is the conclusion. We will handle these proves by introducing first categorical semantics for pattern unification. Before going

Let us start by giving a more formal definition of GB-signatures.

A. Pattern unification as a coequaliser construction

In this section, we assume given a GB-signature $S = (\mathcal{A}, O, \kappa, \alpha)$ and explain how most general unifiers can be thought of as equalisers in a multi-sorted Lawvere theory, as is well-known in the first-order case [19], [2]. We further more provide a formal justification for the error metavariable context \perp .

Lemma 12. *Metavariable contexts and substitutions (with their composition) between them define a category $\text{MCon}(S)$.*

We will show later how this category fully faithfully embeds in a Kleisli category for a monad generated by S .

Remark 13. $\text{MCon}(S)$ is the opposite category of a multi-sorted Lawvere theory: the sorts are the objects of \mathcal{A} . This theory is not freely generated by operations unless \mathcal{A} is discrete. Thus, even the GB-signature induced (as in Example 6) by an empty binding signature is not “free”. When \mathcal{A} is discrete, we recover (multi-sorted) first-order unification.

Since a substitution is precisely a list of terms sharing the same metavariable context Γ , a unification problem for two list of terms is equivalently given by a pair of parallel substitutions

$$\Gamma \xrightarrow[\sigma_2]{\sigma_1} \Delta.$$

Lemma 14. *The most general unifier of two lists of terms $\Delta; n_i \vdash t_i, u_i$, if it exists, is characterised as the coequaliser of \vec{t} as \vec{u} as substitutions from $(N_1 : n_1, \dots)$ to Δ .*

This justifies a common interpretation as (co)equalisers of the two variants of the notation $- \vdash - = - \Rightarrow - \dashv -$ involved in Figure 2.

Pattern unification is often stated as the existence of a coequaliser on the condition that there is a unifier. It turns out that we can get rid of this condition by considering the category $\text{MCon}(S)$ freely extended with a terminal object \perp , as we now explain.

Definition 15. Given a category \mathcal{B} , let \mathcal{B}_\perp denote the category \mathcal{B} extended freely with a terminal object \perp .

Adding a terminal object results in adding a terminal cocone to all diagrams. As a consequence, we have the following lemma.

Lemma 16. *Let J be a diagram in a category \mathcal{B} . The following are equivalent:*

- 1) J has a colimit as long as there exists a cocone;
- 2) J has a colimit in \mathcal{B}_\perp .

The following result is also useful.

Lemma 17. *Given a category \mathcal{B} , the canonical embedding functor $\mathcal{B} \rightarrow \mathcal{B}_\perp$ creates colimits.*

As a consequence,

- 1) whenever the colimit in $\text{MCon}_\perp(S)$ is not \perp , it is also a colimit in $\text{MCon}(S)$;
- 2) existing colimits in $\text{MCon}(S)$ are also colimits in $\text{MCon}_\perp(S)$;
- 3) in particular, coproducts in $\text{MCon}(S)$, which are computed as union of metavariable contexts, are also coproducts in $\text{MCon}_\perp(S)$.

Categorically speaking, our pattern unification algorithm provides an explicit proof of the following statment.

Theorem 18. *Given any pattern-friendly signature S , the category $\text{MCon}_\perp(S)$ has coequalisers.*

B. Generalised binding signatures

Definition 19. A generalised binding signature, or GB-signature, is a tuple $(\mathcal{A}, O, \kappa, \alpha)$ consisting of

- a small category \mathcal{A} of arities and renamings between them;
- a functor $O : \mathcal{A} \rightarrow \text{Set}$ of operations;
- a functor $\kappa : (\int O)^{op} \rightarrow \mathbb{F}$;
- a functor $\alpha : (\int \kappa)^{op} \rightarrow \mathcal{A}$

where

- \mathbb{F} is the category of finite cardinals and any morphisms between them;
- $\int O$ denotes the category of elements of O whose objects are pairs of an object a of \mathcal{A} and an element o in O_a , and morphisms between (a, o) and (a', o') are morphisms $f : a \rightarrow a'$ such that $o\{f\} = o'$;
- $\int \kappa$ denotes the category of elements of $(\int O)^{op} \xrightarrow{\kappa} \mathbb{F} \hookrightarrow \text{Set}$. Objects are triples (a, o, i) , where a is an object of \mathcal{A} , $o \in O_a$, and $j \in \{1, \dots, \kappa_o\}$, and a morphism in $\int \kappa$ between (a, o, i) and (a', o', i') is a morphism $f : a' \rightarrow a$ such that $o = o'\{f\}$ and $i' = \epsilon(o, f)_i$, where $\epsilon(o, f)_-$ is the morphism between κ_o and $\kappa_{o'}$ induced by the functorial action of κ on f .

Any GB-signature $S = (\mathcal{A}, O, \kappa, \alpha)$, generates an endofunctor F_S on $[\mathcal{A}, \text{Set}]$, that we denote by just F when the context is clear, defined by

$$F_S(X)_a = \prod_{o \in O_a} X_{\bar{o}_1} \times \dots \times X_{\bar{o}_{\kappa_o}}.$$

Lemma 20. *F is finitary and generates a free monad that restricts to a monad T . Moreover, TX is the initial algebra of $Z \mapsto X + FZ$.*

Proof. F is finitary because filtered colimits commute with finite limits [14, Theorem IX.2.1] and colimits. The free monad construction is due to [17]. \square

We now introduce our conditions for the correctness criterion.

Definition 21. A GB-signature $S = (\mathcal{A}, O, \kappa, \alpha)$ is said *pattern-friendly* if \mathcal{A} has finite connected limits, all morphisms are monomorphic, and moreover F restricts to an endofunctor on the full subcategory \mathcal{C} of functors from \mathcal{A} to \mathbf{Set} preserving finite limits.

Example 22. When $\mathcal{A} = \mathbb{F}_m$, the category \mathcal{C} is equivalent to the category of nominal sets [8].

Lemma 23. The syntax generated by a GB-signature (see Definition 4) is recovered as free algebras for F . More precisely, given a metavariable context $\Gamma = (M_1 : m_1, \dots, M_p : m_p)$,

$$T(\underline{\Gamma})_a \cong \{t | \Gamma; a \vdash t\}$$

where $\underline{\Gamma} : \mathcal{A} \rightarrow \mathbf{Set}$ is defined as the coproduct of representable functors $\coprod_i ym_i$, mapping a to $\coprod_i \text{hom}_{\mathcal{A}}(m_i, a)$.

Notation 24. Given a metavariable context Γ . We sometimes denote $\underline{\Gamma}$ just by Γ .

If $\Gamma = (M_1 : m_1, \dots, M_p : m_p)$ and Δ are metavariable contexts, a Kleisli morphism $\sigma : \Gamma \rightarrow T\Delta$ is equivalently given (by the Yoneda Lemma and the universal property of coproducts) by a metavariable substitution from Γ to Δ . Moreover, Kleisli composition corresponds to composition of substitutions. This provides a formal link between the category of metavariable contexts $\text{MCon}(S)$ and the Kleisli category of T .

Lemma 25. The category $\text{MCon}(S)$ is equivalent to the full subcategory of Kl_T spanned by coproducts of representable functors.

We will exploit this characterisation to prove various properties of this category when the signature is pattern-friendly.

C. Properties of pattern-friendly signatures

We now assume given a pattern-friendly signature $S = (\mathcal{A}, O, \kappa, \alpha)$.

Lemma 26. \mathcal{C} is closed under limits, coproducts, and filtered colimits. Moreover, it is cocomplete.

Lemma 27. T restricts as a monad on \mathcal{C} .

By right continuity of the homset bifunctor, any representable functor is in \mathcal{C} and thus the embedding $\mathcal{C} \rightarrow [\mathcal{A}, \mathbf{Set}]$ factors the Yoneda embedding $\mathcal{A}^{op} \rightarrow [\mathcal{A}, \mathbf{Set}]$.

Notation 28. We set $\mathcal{D} = \mathcal{A}^{op}$ and denote the fully faithful Yoneda embedding as $\mathcal{D} \xrightarrow{K} \mathcal{C}$. We denote by $\mathcal{D}^+ \xrightarrow{K^+} \mathcal{C}$ the full subcategory of \mathcal{C} consisting of finite coproducts of objects of \mathcal{D} . Given a metavariable context $\Gamma = (M_1 : m_1, \dots, M_p : m_p)$, we denote by $\underline{\Gamma}$ or even Γ the coproduct $\coprod_{i \in \{1, \dots, p\}} Km_i$.

Remark 29. \mathcal{D}^+ is equivalent to the category of finite families of objects of \mathcal{A} , and can be thought of as the category of metavariable contexts and metavariable renamings.

We now abstract the situation by listing a number of properties that we will use to justify the unification algorithm.

Property 30. The following properties hold.

- (i) \mathcal{D} has finite connected colimits.
- (ii) $K : \mathcal{D} \rightarrow \mathcal{C}$ preserves finite connected colimits.
- (iii) Given any morphism $f : a \rightarrow b$ in \mathcal{D} , the morphism Kf is epimorphic.
- (iv) Coproduct injections $A_i \rightarrow \coprod_j A_j$ in \mathcal{C} are monomorphisms.
- (v) For each $d \in \mathcal{D}$, the object Kd is connected, i.e., any morphism $Kd \rightarrow \coprod_i A_i$ factors through exactly one coproduct injection $A_j \rightarrow \coprod_i A_i$.

Proof. (i) We assume that \mathcal{A} has finite connected limits. Hence, its opposite category $\mathcal{D} = \mathcal{A}^{op}$ has finite connected colimits. \square

(ii) Let $y : \mathcal{A}^{op} \rightarrow [\mathcal{A}, \mathbf{Set}]$ denote the Yoneda embedding and $J : \mathcal{C} \rightarrow [\mathcal{A}, \mathbf{Set}]$ denote the canonical embedding, so that

$$y = J \circ K. \quad (3)$$

Now consider a finite connected limit $\lim F$ in \mathcal{A} . Then,

$$\begin{aligned} \mathcal{C}(K \lim F, X) &\cong [\mathcal{A}, \mathbf{Set}](JK \lim F, JX) \\ &\quad (J \text{ is fully faithful}) \\ &\cong [\mathcal{A}, \mathbf{Set}](y \lim F, JX) \quad (\text{By (3)}) \\ &\cong JX(\lim F) \quad (\text{By the Yoneda Lemma.}) \\ &\cong \lim(JX \circ F) \\ &\quad (X \text{ preserves finite connected limits}) \\ &\cong \lim([\mathcal{A}, \mathbf{Set}](yF-, JX)) \\ &\quad (\text{By the Yoneda Lemma}) \\ &\cong \lim([\mathcal{A}, \mathbf{Set}](JKF-, JX)) \quad (\text{By (3)}) \\ &\cong \lim \mathcal{C}(KF-, X) \quad (J \text{ is full and faithful}) \\ &\cong \mathcal{C}(\text{colim } KF, X) \\ &\quad (\text{By left continuity of the hom-set bifunctor}) \end{aligned}$$

Thus, $K \lim F \cong \text{colim } KF$.

(iii) A morphism $f : a \rightarrow b$ is epimorphic if and only if the following square is a pushout [14, Exercise III.4.4]

$$\begin{array}{ccc} a & \xrightarrow{f} & b \\ f \downarrow & & \parallel \\ b & \xlongequal{\quad} & b \end{array}$$

We conclude by (ii), because all morphisms in $\mathcal{D} = \mathcal{A}^{op}$ are epimorphic by assumption.

(iv) This follows from Lemma 26, because a morphism $f : A \rightarrow B$ is monomorphic if and only if the following square is a pullback

$$\begin{array}{ccc}
A & \xlongequal{\quad} & A \\
\parallel & & \downarrow f \\
A & \xrightarrow{f} & B
\end{array}$$

(v) This follows from coproducts being computed pointwise (Lemma 26), and representable functors being connected, by the Yoneda Lemma.

Proof. . \square

Remark 31. Unification of two metavariables as pullbacks or equalisers in \mathcal{A} crucially relies on Property 30.(ii), which holds because we restrict to functors preserving finite connected limits.

Lemma 32. O preserves finite connected limits.

Proof. O is isomorphic to $F(1)$, where 1 is the constant functor mapping everything to the singleton set $\{0\}$. Since 1 trivially preserves limits, it is in \mathcal{C} and thus $F(1) \cong O$ also is. \square

Notation 33. We denote by $\mathcal{L}^+ : \mathcal{D}^+ \rightarrow \text{MCon}(S)$ the functor which is the identity on objects and postcomposes any morphism $A \rightarrow B$ by $\eta_B : B \rightarrow TB$, and by \mathcal{L} the functor $\mathcal{D} \hookrightarrow \mathcal{D}^+ \xrightarrow{\mathcal{L}^+} \text{MCon}(S)$.

Remark 34. Intuitively, the functor \mathcal{L} maps a morphism $x \in \text{hom}_{\mathcal{A}}(m, a)$ to the morphism $ya \rightarrow T(M : m)$ selecting (by the Yoneda Lemma) the term $M(x)$. Thus, a metavariable application $\Gamma; a \vdash M(x)$ corresponds to the composition $\mathcal{L}x[in_M]$, where in_M is the embedding of $M : m$ into Γ .

Property 35. The functor $\mathcal{D} \xrightarrow{\mathcal{L}} \text{MCon}(S)$ preserves finite connected colimits.

V. SOUNDNESS OF THE UNIFICATION PHASE

In this section, we discuss soundness of some of the rules of the main unification phase in Figure 2, which computes a coequaliser in $\text{MCon}_{\perp}(S)$. More specifically, we discuss the rule sequential rule U-SPLIT (Section §V-A), the rule U-FLEXFLEX unifying metavariable with itself (Section §V-B), and the failing rule U-CYCLIC for cyclic unification of a metavariable with a term which includes it deeply (Section §V-C).

A. Sequential unification (rule U-SPLIT)

The rule U-SPLIT follows from a stepwise construction of coequalisers valid in any category, as noted by [19, Theorem 9] when studying first-order unification: if the first two diagrams below are coequalisers, then the last one as well.

$$\begin{array}{ccc}
A_1 \xrightarrow[t_1]{u_1} \Gamma \xrightarrow{\sigma_1} \Delta_1 & & A_2 \xrightarrow[t_2]{u_2} \Gamma \xrightarrow{\sigma_1} \Delta_1 \xrightarrow{\sigma_2} \Delta_2 \\
& & \downarrow u_2 \quad \downarrow \sigma_1 \\
& & \Gamma \xrightarrow{\sigma_1} \Delta_1 \xrightarrow{\sigma_2} \Delta_2
\end{array}$$

$$A_1 + A_2 \xrightarrow[u_1, u_2]{t_1, t_2} \Gamma \xrightarrow{\sigma_2 \circ \sigma_1} \Delta_2$$

B. Flex-Flex, same metavariable (rule U-FLEXFLEX)

Here we detail unification of $M(x)$ and $M(y)$, for $x, y \in \text{hom}_{\mathcal{A}}(m, a)$. By Remark 34, $M(x) = \mathcal{L}x[in_M]$ and $M(y) = \mathcal{L}y[in_M]$. We exploit the following lemma with $u = \mathcal{L}x$ and $v = \mathcal{L}y$.

Lemma 36. In any category, denoting morphism composition $g \circ f$ by $f[g]$, the following rule applies:

$$\frac{B \vdash u = v \Rightarrow h \dashv C}{B + D \dashv u[in_B] = v[in_B] \Rightarrow h + 1_D \dashv C + D}$$

In other words, if the below left diagram is a coequaliser, then so is the below right diagram.

$$\begin{array}{ccc}
A \xrightarrow[u]{v} B \xrightarrow{h} C & & A \xrightarrow[u]{v} B \xrightarrow{in_B} B + D \xrightarrow{h+1_D} C + D \\
& & \downarrow u \quad \downarrow v \\
& & B \xrightarrow{in_B} B + D \xrightarrow{h+1_D} C + D
\end{array}$$

It follows that it is enough to compute the coequaliser of $\mathcal{L}x$ and $\mathcal{L}y$. Furthermore, by Property 30.(i) and Property 35, it can be computed as the image of the coequaliser of x and y , thus justifying the rule U-FLEXFLEX.

C. Flex-rigid, cyclic (rule U-CYCLIC)

The rule U-CYCLIC handles unification of $M(x)$ and a term u such that u is rigid and M occurs in u . In this section, we show in this section that indeed there is no unifier. More precisely, we prove in Corollary 41 below that if there is a unifier, for a term u in and a metavariable application $M(x)$, then either M occurs at top-level in u , or it does not occur at all. The argument follows the basic intuition that $t = u[M \mapsto t]$ is impossible if M occurs deeply in u because the sizes of both hand sides can never match. To make this statement precise, we need some recursive definitions and properties of size.

Definition 37. The size⁴ $|t| \in \mathbb{N}$ of a term t is recursively defined by $|M(x)| = 0$ and $|o(\vec{t})| = 1 + |\vec{t}|$, with $|\vec{t}| = \sum_i t_i$.

We will also need to count the occurrences of a metavariables in a term.

Definition 38. For any term t we define $|t|_M$ recursively by $|M(x)|_M = 1$, $|N(x)|_M = 0$ if $N \neq M$, and $|o(\vec{t})|_M = |\vec{t}|_M$ with the sum convention as above for $|\vec{t}|_M$.

Lemma 39. For any term $\Gamma, M : m; a \vdash t$, if $|t|_M = 0$, then factors $\Gamma; a \vdash t$. Moreover, for any $\Gamma = (M_1 : m_1, \dots, M_n : am_n)$, and well-formed term t in context $\Gamma; a$, and substitution $\sigma : \Gamma \rightarrow \Delta$, we have $|t[\sigma]| = |t| + \sum_i |t|_{M_i} \times |\sigma_i|$.

Corollary 40. For any term t in context $\Gamma, M : m; a$, substitution $\sigma : \Gamma \rightarrow \Delta$, morphism $x \in \text{hom}_{\mathcal{A}}(m, a)$ and u in context $\Delta; u$, we have $|t[\sigma, M \mapsto u]| \geq |t| + |u| \times |t|_M$ and $|M(x)[u]| = |u|$.

⁴The difference with the notion of size introduced in section §III is that metavariables are of size 0.

Corollary 41. *Let t be a term in context $\Gamma, M : m; a$ and $x \in \text{hom}_{\mathcal{A}}(m, a)$. Then, either $t = M(y)$ for some $y \in \text{hom}_{\mathcal{A}}(m, a)$, or $\Gamma; a \vdash t$.*

Proof. Since $t[\sigma, u] = M(x)[u]$, we have $|t[\sigma, u]| = |M(x)[u]|$. Corollary 40 implies $|u| \geq |t| + |u| \times |t|_M$. Therefore, either $|t|_M = 0$ and we conclude by Lemma 39, or $|t|_M = 1$ and $|t| = 0$ and so t is $M(y)$ for some y . \square

VI. SOUNDNESS OF THE NON-CYCLIC PHASE

The non-cyclic phase is concerned with unifying a list of terms $\Gamma; n_i \vdash t_i$ with a list of fresh metavariable applications $M_1(x_1), \dots, M_p(x_p)$, in the extended metavariable context $\Gamma, M_1 : m_1, \dots, M_p : x_p$. Categorically speaking, we are looking at a coequalising diagram in $\text{MCon}(S)$.

$$\begin{array}{ccccc} \overrightarrow{N : n} & & \xrightarrow{\vec{t}} & \Gamma & \xrightarrow{in_1} & \Gamma, \overrightarrow{M : m} \\ & \searrow & & & \nearrow & \\ & \overrightarrow{M(x)} & \xrightarrow{\vec{M}} & \overrightarrow{M : m} & \xrightarrow{in_2} & \Gamma, \overrightarrow{M : m} \end{array}$$

The P-SPLIT rule is a straightforward adaption of the U-SPLIT rule specialised to those specific coequaliser diagrams.

Remark 42. A unifier $\Gamma, \overrightarrow{M : m} \rightarrow \Delta$ splits into two components: a substitution $\sigma : \Gamma \rightarrow \Delta$ and a substitution \vec{u} from $\overrightarrow{N : n}$ to $\overrightarrow{M : m}$ such that $t_i[\sigma] = u_i\{x_i\}$ for each $i \in \{1, \dots, p\}$

The coequaliser $\sigma, \vec{u} : (\Gamma, \overrightarrow{M : m}) \rightarrow \Delta$ is equivalently characterised as a pushout

$$\begin{array}{ccc} \overrightarrow{N : n} & \xrightarrow{\vec{M(x)}} & \overrightarrow{M : m} \\ \downarrow \vec{t} & & \downarrow \vec{u} \\ \Gamma & \xrightarrow{\sigma} & \Delta \end{array}$$

This justifies a common interpretation as pushouts of the two variants of the notation $\vdash - :> - \Rightarrow -; -$ involved in Figure 2, in \mathcal{A}^{op} and $\text{MCon}(S)$.

In the following sections, we detail soundness of the rules for the rigid case (Section §VI-A) and then for the flex case (Section §VI-B).

A. Rigid (rules P-RIG and P-FAIL)

In this section, we describe soundness for the rules P-RIG and P-FAIL which handle the non-cyclic unification of $M(x)$ with $\Gamma; a \vdash o(\vec{t})$ in the metavariable context $\Gamma, M : m$. By Remark 42, a unifier is a substitution $\sigma : \Gamma \rightarrow \Delta$ and a term u such that

$$o(\vec{t}[\sigma]) = u\{x\}. \quad (4)$$

Now, u is either some $M(y)$ or $o'(\vec{v})$. But in the first case, $u\{x\} = M(y)\{x\} = M(x \circ y)$, contradicting Equation (4). Therefore, $u = o'(\vec{v})$ for some $o' \in O_m$ and $\vec{v} = (v_1, \dots, v_{\kappa_{o'}})$ such that $\Delta; \vec{o}'_i \vdash v_i$. Then, $u\{x\} = (o'\{x\})(v_{\epsilon(o', x)_1}\{x_1^{o'}\}, \dots)$. It follows from Equation (4) that $o = o'\{x\}$, and $t_i[\sigma] = v_{\epsilon(o', x)_i}\{x_i^{o'}\}$.

Note that there is at most one o' such that $o = o'\{x\}$, by Property 30.(iii). In this case, a unifier is equivalently given by a substitution $\sigma : \Gamma \rightarrow \Delta$ and a list of terms $\vec{v} = (v_1, \dots, v_{\kappa_{o'}})$ such that $\Delta; \vec{o}'_i \vdash v_i$ and $t_i[\sigma] = v_{\epsilon(o', x)_i}\{x_i^{o'}\}$ for all $i \in \{1, \dots, p\}$. But, by Remark 42, this is precisely the data for a unifier of \vec{t} and $M_1(x_1^{o'}), \dots, M_p(x_p^{o'})$. This actually induces an isomorphism between the two categories of unifiers, thus justifying the rules P-RIG and P-FAIL.

B. Flex (rule P-FLEX)

The rule P-FLEX handles unification of $\Gamma, N : n; a \vdash N(x)$ and $M(y)$ where M is fresh.

Note that $M(y) = \mathcal{L}y$ as a substitution $(A : a) \rightarrow (M : m)$ while $N(x) = \mathcal{L}x[in_N]$, by Remark 34. Thanks to the following lemma, it is actually enough to compute the pushout of $\mathcal{L}x$ and $\mathcal{L}y$.

Lemma 43. *In any category, denoting morphism composition by $f \circ g = g[f]$, the following rule applies*

$$\frac{X \vdash g :> f \Rightarrow u; \sigma \vdash Z}{X + Y \vdash g[in_1] :> f \Rightarrow u[in_1]; \sigma + Y \vdash Z + Y}$$

In other words, if the diagram below left is a pushout, then so is the right one.

$$\begin{array}{ccc} A & \xrightarrow{f} & B \\ g \downarrow & & \downarrow u \\ X & \xrightarrow{\sigma} & Z \end{array} \quad \begin{array}{ccc} A & \xrightarrow{f} & B \\ g \downarrow & & \downarrow u \\ X & \xrightarrow{in_1} & Z \\ X + Y & \xrightarrow{\sigma + Y} & Z + Y \end{array}$$

By Property 30.(i) and Property 35, the pushout of $\mathcal{L}x$ and $\mathcal{L}y$ is the image by \mathcal{L} of the pushout of x and y , thus justifying the rule P-FLEX.

VII. APPLICATIONS

In the examples, we motivate the definition of the category \mathcal{A} based on what we expect from metavariable arities, following Remark ??.

In all our examples, O is a coproduct $\coprod_{\ell \in V} O_\ell$ for some $O_\ell : \mathcal{A} \rightarrow \text{Set}$ and $J_{a, in_\ell(o)} = \gamma_\ell$ for some finite cardinal γ_ℓ . In this case, L is equivalently given by a family of functors $H_{\ell, j} : \int O_\ell \rightarrow \text{Set}$ for each $j \in \gamma_\ell$ by $L_{a, in_\ell(o), j} = H_{\ell, j}(a, o)$. Then,

$$F(X)_a \cong \prod_{\ell \in V} \prod_{o \in O_\ell} \prod_{j \in \gamma_\ell} X_{H_{\ell, j}(a, o)}$$

A. Simply-typed second-order syntax

In this section, we present the example of simply-typed λ -calculus. Our treatment generalises to any second-order binding signature (see [6]).

Let T denote the set of simple types generated by a set of simple types. A metavariable arity $\tau_1, \dots, \tau_n \vdash \tau_f$ is given by a list of input types τ_1, \dots, τ_n and an output type τ_f . Substituting a metavariable $M : (\Gamma \vdash \tau)$ with another $M' : (\Gamma' \vdash \tau')$ requires that $\tau = \tau'$ and involves an injective

renaming $\Gamma \rightarrow \Gamma'$. Thus, we consider $\mathcal{A} = \mathbb{F}_m[T] \times T$, where $\mathbb{F}_m[T]$ is the category of finite lists of elements of T and injective renamings between them.

Table I summarises the definition of the endofunctor F on $[\mathcal{A}, \text{Set}]$ specifying the syntax, where $|\Gamma|_\tau$ denotes the number (as a cardinal set) of occurrences of τ in Γ .

B. Arguments as sets

If we think of the arguments of a metavariable as specifying the available variables, then it makes sense to assemble them in a set rather than in a list. This motivates considering the category $\mathcal{A} = \mathbb{I}$ whose objects are natural numbers and a morphism $n \rightarrow p$ is a subset of $\{0, \dots, p-1\}$ of cardinal n . For instance, \mathbb{I} can be taken as subcategory of \mathbb{F}_m consisting of strictly increasing injections, or as the subcategory of the augmented simplex category consisting of injective functions. Again, we can define the endofunctor for λ -calculus as in Section §???. Then, a metavariable takes as argument a set of available variables, rather than a list of distinct variables. In this approach, unifying two metavariables (see the rules U-FLEXFLEX and P-FLEX) amount to computing a set intersection.

C. Linear ordered λ -calculus

Properties of Terms in Continuation-Passing Style in an Ordered Logical Framework

We take the example of linear λ -calculus. A metavariable arity is then a natural number specifying the number of arguments, and because of linearity, a metavariable can only be substituted with another one with the same arity. We therefore $\mathcal{A} = \mathcal{P}$ to be category whose objects are finite cardinals and morphisms are bijections as in [20].

Each typing rule provides an operation

$$\begin{array}{lcl} \text{Variable} & 1 \vdash * & y_1 \\ \text{Application} & \frac{n \vdash t \quad m \vdash u}{n+m \vdash t u} & X_n \times X_m \times y_{n+m} \end{array}$$

D. Intrinsic polymorphic syntax

We present intrinsic system F , following [12]. Let $S : \mathbb{F}_m \rightarrow \text{Set}$ mapping n to the set S_n of types for system F taking free type variables in $\{0, \dots, n-1\}$. Intuitively, a metavariable arity $n; \sigma_1, \dots, \sigma_p \vdash \tau$ specifies the number n of free type variables, the list of input types $\vec{\sigma}$, and the output type τ , all living in S_n . Substituting a metavariable $M : (n; \vec{\sigma} \vdash \tau)$ with another $M' : (n'; \vec{\sigma}' \vdash \tau')$ requires a choice $(\alpha_0, \dots, \alpha_{n-1})$ of n distinct type variables among $\{0, \dots, n'-1\}$, such that $\tau[\vec{\alpha}] = \tau'$, and an injective renaming $\vec{\sigma}'[\vec{\alpha}] \rightarrow \vec{\sigma}$. We therefore consider the category \mathcal{A} of metavariable arities where a morphism between $n; \Gamma \vdash \tau$ and $n'; \Gamma' \vdash \tau'$ is a morphism $\sigma : n \rightarrow n'$ in \mathbb{F}_m such that $\tau[\sigma] = \tau'$, and a renaming $\Gamma[\sigma] \rightarrow \Gamma'$. More formally, \mathcal{A} is the op-lax colimit of $n \mapsto \mathbb{F}_m[S_n] \times S_n$. The intrinsic syntax of system F can then be specified as in Table I.

To understand how unification of two metavariables works (see the rules U-RIGRIG and P-RIG), let us explain how finite connected limits are computed in \mathcal{A} .

Let us introduce the category \mathcal{A}' whose definition follows that of \mathcal{A} , but without the output types: objects are pairs of a natural number n and an element of S_n . Note that this is op-lax colimit of $n \mapsto \mathbb{F}_m[S_n]$, and there is an obvious projection $\mathcal{A} \rightarrow \mathcal{A}'$, which creates finite limits, as we will show.

Lemma 44. *\mathcal{A}' has finite limits, and the projection functor $\mathcal{A} \rightarrow \mathcal{A}'$ preserves them.*

Proof. The crucial point is that \mathcal{A}' is not only op-fibred over \mathbb{F}_m by construction, it is also fibred over \mathbb{F}_m . Intuitively, if $\Gamma \in \mathbb{F}_m[S_n]$ and $f : n' \rightarrow n$ is a morphism in \mathbb{F}_m , then $f_! \Gamma \in \mathbb{F}_m[S_{n'}]$ is essentially Γ restricted to elements of S_n that are in the image of S_f . Note that $f_!$ is right adjoint to $\Gamma \mapsto \Gamma[f]$, and is thus continuous. We now apply [11, Theorem 4.2 and Proposition 4.1]: each $\mathbb{F}_m[S_n]$ has those limits. \square

Lemma 45. *The projection functor $\mathcal{A} \rightarrow \mathcal{A}'$ creates finite limits.*

Proof. Let $d : I \rightarrow \mathcal{A}$ be a functor. We denote d_i by $n_i; \Gamma_i \vdash \tau_i$. Let $n; \Gamma$ be the limit of $i \mapsto n_i; \Gamma_i$ in \mathcal{A}' . By the previous lemma, n is the limit of $i \mapsto n_i$. Note that $S : \mathbb{F}_m \rightarrow \text{Set}$ preserves finite connected limits. Thus, we can define $\tau \in S_n$ as corresponding to the universal function $1 \rightarrow S_n$ factorising the cone $(1 \xrightarrow{\tau_i} S_{n_i})_i$.

It is easy to check that $n; \Gamma \vdash \tau$ is the limit of d . \square

More concretely, a finite connected limit of $i \mapsto n_i; \Gamma_i \vdash \tau_i$ in \mathcal{A} is computed as follows:

- 1) compute the limit n of $(n_i)_i$ in \mathbb{F}_m , denoting $p_i : n \rightarrow n_i$ is the canonical projections;
- 2) define τ as the (only) element of S_n such that $\tau[p_i] = \tau_i$
- 3) define Γ as the limit in $\mathbb{F}_m[S_n]$ of $p_{i,!} \Gamma_i$, where $p_{i,!} : \mathbb{F}_m[S_{n_i}] \rightarrow \mathbb{F}_m[S_n]$ is the reindexing functor described in the proof of Lemma 44.

This means that to unify $M(\vec{\alpha}; \vec{\tau})$ with $M(\vec{\alpha}'; \vec{\tau}')$, with M of arity p ; $\vec{u} \vdash u'$, we first need to compute the vector of common position \vec{i} between $\vec{\alpha}$ and $\vec{\alpha}'$, i.e, the largest vector $(i_1 < \dots < i_n)$ such that $\alpha_{i_j} = \alpha'_{i_j}$. Then, we consider $\vec{\sigma}$ and $\vec{\sigma}'$ so that $\vec{\sigma}[j \mapsto i_{j+1}]$ is the sub-list of $\vec{\tau}$ that only use type variables in $\alpha_{i_1}, \dots, \alpha_{i_n}$, and similarly for $\vec{\sigma}'$ and $\vec{\tau}'$. Finally, we define (t_1, \dots, t_m) as the vector of common positions between $\vec{\sigma}$ and $\vec{\sigma}'$. The most general unifier is then $M \mapsto N(\vec{i}, \vec{t})$ for a fresh metavariable N of arity $n; \sigma_{t_1}, \dots, \sigma_{t_m} \vdash t'$, where $t' = u'[i_{j+1} \mapsto j]$.

REFERENCES

- [1] Altenkirch, T., Morris, P.: Indexed containers. In: Proceedings of the 24th Annual IEEE Symposium on Logic in Computer Science, LICS 2009, 11-14 August 2009, Los Angeles, CA, USA. pp. 277–285. IEEE Computer Society (2009). <https://doi.org/10.1109/LICS.2009.33>, <https://doi.org/10.1109/LICS.2009.33>
- [2] Barr, M., Wells, C.: Category Theory for Computing Science. Prentice-Hall, Inc., USA (1990)
- [3] D., P.G.: A note on inductive generalization. Machine Intelligence **5**, 153–163 (1970), <https://cir.nii.ac.jp/crid/1573387448853680384>

TABLE I
EXAMPLES OF ENDOFUNCTORS FOR SYNTAX

$$F(X)_a \cong \coprod_{\ell \in V} \coprod_{o \in O_\ell} \coprod_{j \in \gamma_\ell} X_{H_{\ell,j}(a,o)}$$

Simply-typed λ -calculus (Section §VII-A)	Typing rule	$O_\ell(\Gamma \vdash \tau)$	$H_{\ell,j}(\Gamma \vdash \tau)$
	$\frac{x : \tau \in \Gamma}{\Gamma \vdash x : \tau}$	$\{v_i i \in \Gamma _\tau\}$	-
	$\frac{\Gamma \vdash t : \tau' \Rightarrow \tau \quad \Gamma \vdash u : \tau'}{\Gamma \vdash t u : \tau}$	$\{a_{\tau'} \tau' \in T\}$	$H_{-,0} = \Gamma \vdash \tau' \Rightarrow \tau$ $H_{-,1} = \Gamma \vdash \tau'$
	$\frac{\Gamma, x : \tau_1 \vdash t : \tau_2}{\Gamma \vdash \lambda x. t : \tau_1 \Rightarrow \tau_2}$	$\{l_{\tau_1, \tau_2} \tau = (\tau_1 \Rightarrow \tau_2)\}$	$H_{-,0} = \Gamma, \tau_1 \vdash \tau_2$
System F (Section §VII-D)	Typing rule	$O_\ell(n; \Gamma \vdash \tau)$	$H_{\ell,j}(n; \Gamma \vdash \tau, o)$
	$\frac{x : \tau \in \Gamma}{n; \Gamma \vdash x : \tau}$	$\{v_i i \in \Gamma _\tau\}$	-
	$\frac{n; \Gamma \vdash t : \tau' \Rightarrow \tau \quad n; \Gamma \vdash u : \tau'}{n; \Gamma \vdash t u : \tau}$	$\{a_{\tau'} \tau' \in S_n\}$	$H_{-,0} = n; \Gamma \vdash \tau' \Rightarrow \tau$ $H_{-,1} = n; \Gamma \vdash \tau'$
	$\frac{n; \Gamma, x : \tau_1 \vdash t : \tau_2}{n; \Gamma \vdash \lambda x. t : \tau_1 \Rightarrow \tau_2}$	$\{l_{\tau_1, \tau_2} \tau = (\tau_1 \Rightarrow \tau_2)\}$	$H_{-,0} = n; \Gamma, \tau_1 \vdash \tau_2$
	$\frac{n; \Gamma \vdash t : \forall \tau_1 \quad \tau_2 \in S_n}{n; \Gamma \vdash t : \tau_2 : \tau_1[\tau_2]}$	$\{A_{\tau_1, \tau_2} \tau = \tau_1[\tau_2]\}$	$H_{-,0} = n; \Gamma \vdash \forall \tau_1$
	$\frac{n+1; wk(\Gamma) \vdash t : \tau}{n; \Gamma \vdash \Lambda t : \forall \tau}$	$\{\Lambda_{\tau'} \tau = \forall \tau'\}$	$H_{-,0} = n+1; wk(\Gamma) \vdash \tau'$

- [4] De Bruijn, N.G.: Lambda-calculus notation with nameless dummies, a tool for automatic formula manipulation, with application to the church-rosser theorem. *Indagationes Mathematicae* **34**, 381–392 (1972)
- [5] Dunfield, J., Krishnaswami, N.R.: Sound and complete bidirectional typechecking for higher-rank polymorphism with existentials and indexed types. *Proc. ACM Program. Lang.* **3**(POPL), 9:1–9:28 (2019). <https://doi.org/10.1145/3290322>, <https://doi.org/10.1145/3290322>
- [6] Fiore, M.P., Hur, C.K.: Second-order equational logic. In: *Proceedings of the 19th EACSL Annual Conference on Computer Science Logic (CSL 2010)* (2010)
- [7] Fiore, M., Plotkin, G., Turi, D.: Abstract syntax and variable binding. In: *Proc. 14th Symposium on Logic in Computer Science IEEE* (1999)
- [8] Gabbay, M.J., Pitts, A.M.: A new approach to abstract syntax involving binders. In: *Proc. 14th Symposium on Logic in Computer Science IEEE* (1999)
- [9] Goguen, J.A.: What is unification? - a categorical view of substitution, equation and solution. In: *Resolution of Equations in Algebraic Structures, Volume 1: Algebraic Techniques*. pp. 217–261. Academic (1989)
- [10] Goldfarb, W.D.: The undecidability of the second-order unification problem. *Theoretical Computer Science* **13**(2), 225–230 (1981). [https://doi.org/https://doi.org/10.1016/0304-3975\(81\)90040-2](https://doi.org/https://doi.org/10.1016/0304-3975(81)90040-2), <https://www.sciencedirect.com/science/article/pii/0304397581900402>
- [11] Gray, J.W.: Fibrated and cofibred categories. In: Eilenberg, S., Harrison, D.K., MacLane, S., Röhrli, H. (eds.) *Proceedings of the Conference on Categorical Algebra*. pp. 21–83. Springer Berlin Heidelberg, Berlin, Heidelberg (1966)
- [12] Hamana, M.: Polymorphic abstract syntax via grothendieck construction (2011)
- [13] Huet, G.P.: A unification algorithm for typed lambda-calculus. *Theor. Comput. Sci.* **1**(1), 27–57 (1975). [https://doi.org/10.1016/0304-3975\(75\)90011-0](https://doi.org/10.1016/0304-3975(75)90011-0), [https://doi.org/10.1016/0304-3975\(75\)90011-0](https://doi.org/10.1016/0304-3975(75)90011-0)
- [14] Mac Lane, S.: *Categories for the Working Mathematician*. No. 5 in Graduate Texts in Mathematics, Springer, 2nd edn. (1998)
- [15] Miller, D.: A logic programming language with lambda-abstraction, function variables, and simple unification. *J. Log. Comput.* **1**(4), 497–536 (1991). <https://doi.org/10.1093/logcom/1.4.497>, <https://doi.org/10.1093/logcom/1.4.497>
- [16] Plotkin, G.: An illative theory of relations. In: Cooper, R., et al. (eds.) *Situation Theory and its Applications*. pp. 133–146. No. 22 in CSLI Lecture Notes, Stanford University (1990)
- [17] Reiterman, J.: A left adjoint construction related to free triples. *Journal of Pure and Applied Algebra* **10**, 57–71 (1977)
- [18] Robinson, J.A.: A machine-oriented logic based on the resolution principle. *J. ACM* **12**(1), 23–41 (jan 1965). <https://doi.org/10.1145/321250.321253>, <https://doi.org/10.1145/321250.321253>
- [19] Rydeheard, D.E., Burstall, R.M.: *Computational category theory*. Prentice Hall International Series in Computer Science, Prentice Hall (1988)
- [20] Tanaka, M.: Abstract syntax and variable binding for linear binders. In: *MFCS '00. LNCS, vol. 1893*. Springer (2000)
- [21] Urban, C., Pitts, A., Gabbay, M.: Nominal unification. In: Baaz, M., Makowsky, J.A. (eds.) *Computer Science Logic*. pp. 513–527. Springer Berlin Heidelberg, Berlin, Heidelberg (2003)
- [22] Vezzosi, A., Abel, A.: A categorical perspective on pattern unification. *RISC-Linz* p. 69 (2014)
- [23] Zhao, J., d. S. Oliveira, B.C., Schrijvers, T.: A mechanical formalization of higher-ranked polymorphic type inference. *Proc. ACM Program. Lang.* **3**(ICFP), 112:1–112:29 (2019). <https://doi.org/10.1145/3341716>, <https://doi.org/10.1145/3341716>

APPENDIX PROOF OF PROPERTY 30

(i) We assume that \mathcal{A} has finite connected limits. Hence, its opposite category $\mathcal{D} = \mathcal{A}^{op}$ has finite connected colimits.

(ii) Let $y : \mathcal{A}^{op} \rightarrow [\mathcal{A}, \text{Set}]$ denote the Yoneda embedding and $J : \mathcal{C} \rightarrow [\mathcal{A}, \text{Set}]$ denote the canonical embedding, so that

$$y = J \circ K. \quad (5)$$

Now consider a finite connected limit $\lim F$ in \mathcal{A} . Then,

$$\begin{aligned}
\mathcal{C}(K \lim F, X) &\cong [\mathcal{A}, \text{Set}](JK \lim F, JX) && (J \text{ is fully faithful}) \\
&\cong [\mathcal{A}, \text{Set}](y \lim F, JX) && (\text{By (5)}) \\
&\cong JX(\lim F) && (\text{By the Yoneda Lemma.}) \\
&\cong \lim(JX \circ F) && (X \text{ preserves finite connected limits}) \\
&\cong \lim([\mathcal{A}, \text{Set}](yF-, JX)) && (\text{By the Yoneda Lemma}) \\
&\cong \lim([\mathcal{A}, \text{Set}](JKF-, JX)) && (\text{By (5)}) \\
&\cong \lim \mathcal{C}(KF-, X) && (J \text{ is full and faithful}) \\
&\cong \mathcal{C}(\text{colim } KF, X) && (\text{By left continuity of the hom-set bifunctor})
\end{aligned}$$

Thus, $K \lim F \cong \text{colim } KF$.

(iii) A morphism $f : a \rightarrow b$ is epimorphic if and only if the following square is a pushout [14, Exercise III.4.4]

$$\begin{array}{ccc}
a & \xrightarrow{f} & b \\
f \downarrow & & \parallel \\
b & \xlongequal{\quad} & b
\end{array}$$

We conclude by (ii), because all morphisms in $\mathcal{D} = \mathcal{A}^{op}$ are epimorphic by assumption.

(iv) This follows from Lemma 26, because a morphism $f : A \rightarrow B$ is monomorphic if and only if the following square is a pullback

$$\begin{array}{ccc}
A & \xlongequal{\quad} & A \\
\parallel & & \downarrow f \\
A & \xrightarrow{f} & B
\end{array}$$

(v) This follows from coproducts being computed pointwise (Lemma 26), and representable functors being connected, by the Yoneda Lemma.