

Generic pattern unification

A categorical approach

Ambroise Lafont Neel Krishnaswami

University of Cambridge

October 2022

What is unification?

$$\underbrace{t} \stackrel{?}{=} \underbrace{u}$$

terms with metavariables M, N, \dots

Unifier = metavariable substitution σ s.t.

$$t[\sigma] = u[\sigma]$$

Most general unifier = unifier σ that uniquely factors any other

$$\forall \delta, \quad t[\delta] = u[\delta] \quad \Leftrightarrow \quad \exists! \delta'. \quad \delta = \delta' \circ \sigma$$

Goal of unification = find the most general unifier

Where is unification used?

First-order unification

No metavariable argument

Examples

- Logic programming (Prolog)
- ML type inference systems

$$(M \rightarrow N) \stackrel{?}{=} (\mathbb{N} \rightarrow M)$$

Second-order unification

$M(\dots)$

Example

- Type theory, proof assistants

$$(\forall x. M(x, u)) \stackrel{?}{=} t$$

Indecidable

Pattern unification [Miller '91]

A **decidable** fragment of second-order unification.

Pattern restriction:

$$M(\underbrace{x_1, \dots, x_n}_{\text{distinct variables}})$$

\exists unification algorithm [Miller '91]

- fails if no unifier
- returns the most general unifier

This work

A **generic** algorithm for pattern unification

- Parameterised by a *signature*
- Categorical semantics

Examples

- *binding signatures*
- Linear syntax (e.g., quantum λ -calculus)
- Intrinsic system F

Related work: algebraic accounts of unification

First-order unification

- Lattice theory [Plotkin '70]
- Category theory
 - [Rydeheard-Burstall '88]
 - [Goguen '89]

Pattern unification

- Category theory
 - [Vezzosi-Abel '14]
normalised λ -terms
 - [This work](#)

Outline

- 1 Pattern unification for pure λ -calculus
 - Syntax
 - Unification algorithm
- 2 Generalisation to binding signatures
- 3 Categorical generalisation
 - A case study: syntax of pure λ -calculus
 - Generic pattern unification
- 4 Example: System F
- 5 Summary of the generic unification algorithm

Outline

- 1 Pattern unification for pure λ -calculus
 - Syntax
 - Unification algorithm
- 2 Generalisation to binding signatures
- 3 Categorical generalisation
 - A case study: syntax of pure λ -calculus
 - Generic pattern unification
- 4 Example: System F
- 5 Summary of the generic unification algorithm

Outline

- 1 Pattern unification for pure λ -calculus
 - Syntax
 - Unification algorithm
- 2 Generalisation to binding signatures
- 3 Categorical generalisation
 - A case study: syntax of pure λ -calculus
 - Generic pattern unification
- 4 Example: System F
- 5 Summary of the generic unification algorithm

Syntax (De Bruijn levels)

Metavariable context $(M_1 : n_1, \dots)$

$$\underbrace{\Gamma; n}_{\text{Variable context}} \vdash t$$

$$\frac{x < n}{\Gamma; n \vdash x} \text{VAR}$$

$$\frac{\Gamma; n \vdash t \quad \Gamma; n \vdash u}{\Gamma; n \vdash t u} \text{APP}$$

$$\frac{\Gamma; n+1 \vdash t}{\Gamma; n \vdash \lambda t} \text{ABS}$$

$$\frac{(M : n) \in \Gamma \quad x_1, \dots, x_n < n \quad x_1, \dots, x_n \text{ distinct}}{\Gamma; n \vdash M(x_1, \dots, x_n)} \text{FLEX}$$

Metavariable substitution

Substitution σ from $\overbrace{(M_1 : m_1, \dots, M_p : m_p)}^{\Gamma}$ to Δ :

$$(\sigma_1, \dots, \sigma_p) \quad \text{s.t.} \quad \Delta; m_i \vdash \sigma_i$$


σ extends to terms:

$$\Gamma; n \vdash t \quad \mapsto \quad \Delta; n \vdash t[\sigma]$$

Outline

- 1 Pattern unification for pure λ -calculus
 - Syntax
 - Unification algorithm
- 2 Generalisation to binding signatures
- 3 Categorical generalisation
 - A case study: syntax of pure λ -calculus
 - Generic pattern unification
- 4 Example: System F
- 5 Summary of the generic unification algorithm

Unification algorithm



Examples


$$\Gamma, M : 2 \vdash M(x, y) = x \Rightarrow (M \mapsto 0) \dashv \vdash \Gamma$$

$$\Gamma, M : 2 \vdash M(x, y) = y \Rightarrow (M \mapsto 1) \dashv \vdash \Gamma$$

$$\frac{\Gamma \vdash t = u \Rightarrow \sigma \dashv \vdash \Delta}{\Gamma \vdash \lambda t = \lambda u \Rightarrow \sigma \dashv \vdash \Delta}$$

$$\frac{\Gamma \vdash "t_1, t_2 = u_1, u_2" \Rightarrow \sigma \dashv \vdash \Delta}{\Gamma \vdash t_1 \ t_2 = u_1 \ u_2 \Rightarrow \sigma \dashv \vdash \Delta}$$

Unifying lists of terms



Examples (lists)

$$\Gamma \vdash () = () \Rightarrow id_{\Gamma} \dashv \Gamma$$

$$\frac{\Gamma \vdash t_1 = u_1 \Rightarrow \sigma_1 \dashv \Delta_1 \quad \Delta_1 \vdash \vec{t}_2[\sigma_1] = \vec{u}_2[\sigma_1] \Rightarrow \sigma_2 \dashv \Delta_2}{\Gamma \vdash t_1, \vec{t}_2 = u_1, \vec{u}_2 \Rightarrow \sigma_1[\sigma_2] \dashv \Delta_2} \text{U-SPLIT}$$

Impossible cases

$$\Gamma \vdash \lambda t = u_1 u_2 \Rightarrow ! \vdash \perp$$

formal "error" context

formal "error" substitution

Unifying a metavariable $M(\vec{x}) \stackrel{?}{=} \dots$

Three cases

- ❶ $M(\vec{x}) \stackrel{?}{=} M(\vec{y})$
- ❷ $M(\vec{x}) \stackrel{?}{=} \dots M(\vec{y}) \dots$
- ❸ $M(\vec{x}) \stackrel{?}{=} u$ and $M \notin u$

Unifying a metavariable with itself

$$M(\vec{x}) \stackrel{?}{=} M(\vec{y})$$

Most general unifier: $M \mapsto M'(\vec{z})$

- \vec{z} = vector of common positions: $x_{\vec{z}} = y_{\vec{z}}$

Formally,

$$\frac{"n \vdash \vec{x} = \vec{y} \Rightarrow \vec{z} \vdash p"}{\Gamma, M : n \vdash M(\vec{x}) = M(\vec{y}) \Rightarrow M \mapsto M'(\vec{z}) \vdash \Gamma, M' : p}$$

Cyclic case

$$M(\vec{x}) \stackrel{?}{=} \dots M(\vec{y}) \dots$$

No unifier (sizes cannot match after substitution)

$$\Gamma \vdash \underbrace{M(\vec{x})} = \underbrace{\dots M(\vec{y}) \dots} \Rightarrow ! \vdash \perp$$

sizes cannot match after substitution

Non cyclic case

$$M(\vec{x}) \stackrel{?}{=} u \ (M \notin u)$$

Most general unifier: $M \mapsto u[\vec{x}^{-1}]$

- Requires

$$fv(u) \subset \vec{x} \tag{1}$$

\Rightarrow **Pruning phase:** enforces (1) by restricting metavariable arities.

Example

$$\begin{array}{ccc} M(x) & \stackrel{?}{=} & N(x, y) \\ N(x, y) & \xrightarrow{\text{pruning}} & N'(x) \end{array}$$

Pruning phase

$$\begin{array}{c}
 \text{\textit{u} after pruning and renamed by } \vec{x}^{-1} \\
 \Gamma \vdash u :> M(\vec{x}) \Rightarrow \overbrace{v}^{}; \sigma \vdash \Delta \\
 \begin{array}{c} \nearrow \\ M \notin \Gamma, u \end{array} \qquad \begin{array}{c} \nearrow \\ \sigma : \Gamma \rightarrow \Delta \\ \text{pruning substitution} \end{array}
 \end{array}$$

Intuition: $\sigma, M \mapsto v = \text{most general unifier for } u \stackrel{?}{=} M(\vec{x})$

‘v replaces M’

Pruning a metavariable

$$M(\vec{x}) \stackrel{?}{=} N(\vec{y})$$

Most general unifier: $M \mapsto N'(\vec{l})$, $N \mapsto N'(\vec{r})$ such that

$$x_{\vec{l}} = y_{\vec{r}}$$

$$\frac{"n \vdash \vec{x} :> \vec{y} \Rightarrow \vec{l}; \vec{r} \vdash p"}{\Gamma, N : n \vdash N(\vec{x}) :> M(\vec{y}) \Rightarrow N'(\vec{l}); N \mapsto N'(\vec{r}) \vdash \Gamma, N' : p}$$

Pruning: other examples

$$\frac{}{\Gamma \vdash x_i :> M(x_0, \dots, x_n) \Rightarrow i; id_{\Gamma} \dashv \Gamma} \qquad \frac{y \notin \vec{x}}{\Gamma \vdash y :> M(\vec{x}) \Rightarrow !; ! \dashv \perp}$$

$$\frac{\Gamma \vdash t :> M_1(\vec{x}, \overbrace{n}^{\text{bound variable}}) \Rightarrow v; \sigma \dashv \Delta}{\Gamma \vdash \lambda t :> M(\vec{x}) \Rightarrow \lambda v; \sigma \dashv \Delta}$$

$$\frac{"\Gamma \vdash t, u :> M_1(\vec{x}), M_2(\vec{x}) \Rightarrow v_1, v_2; \sigma \dashv \Delta"}{\Gamma \vdash t \ u :> M(\vec{x}) \Rightarrow v_1 \ v_2; \sigma \dashv \Delta}$$

Pruning multi-terms

$$\begin{array}{c}
 \Gamma \vdash u_1, \dots, u_n :> M_1(\vec{x}_1), \dots, M_n(\vec{x}_n) \Rightarrow v_1, \dots, v_n; \sigma \dashv \Delta \\
 \begin{array}{ccc}
 \nwarrow & \nearrow & \nwarrow \\
 \Gamma; (n_i) \vdash u_i & & \Delta; m_i \vdash v_i \\
 \nearrow & & \nearrow \\
 (M_i : m_i) \notin \Gamma, u_j & & \sigma : \Gamma \rightarrow \Delta
 \end{array}
 \end{array}$$

Outline

- 1 Pattern unification for pure λ -calculus
 - Syntax
 - Unification algorithm
- 2 Generalisation to binding signatures
- 3 Categorical generalisation
 - A case study: syntax of pure λ -calculus
 - Generic pattern unification
- 4 Example: System F
- 5 Summary of the generic unification algorithm

Parameterisation by a *signature*

Binding signature for pure λ -calculus

 $app : (0, 0)$ $abs : (1)$

number of bound variables in the argument

Example: pruning an operation

$$o : (\alpha_1, \dots, \alpha_p)$$

$$\frac{\Gamma \vdash \vec{t} :> M_1(\vec{x}, \overbrace{n, \dots, n + \alpha_1 - 1}^{\text{bound variables}}), \dots, M_p(\dots) \Rightarrow \vec{u}; \sigma \vdash \Delta}{\Gamma \vdash o(\vec{t}) :> N(\vec{x}) \Rightarrow o(\vec{u}); \sigma \vdash \Delta}$$

Semantics

$$\begin{array}{ll} \Gamma \vdash t = u \Rightarrow \sigma \dashv \Delta & \Leftrightarrow \quad \sigma = \text{coequaliser of } t \text{ and } u \\ \Gamma \vdash t :> f \Rightarrow v; \sigma \dashv \Delta & \Leftrightarrow \quad \sigma, v = \text{pushout of } t \text{ and } f \end{array}$$

... in the *second-order algebraic theory*¹ of S (cf next section)

¹[Fiore-Mahmoud 2010]

Outline

- 1 Pattern unification for pure λ -calculus
 - Syntax
 - Unification algorithm
- 2 Generalisation to binding signatures
- 3 Categorical generalisation
 - A case study: syntax of pure λ -calculus
 - Generic pattern unification
- 4 Example: System F
- 5 Summary of the generic unification algorithm

Outline

- 1 Pattern unification for pure λ -calculus
 - Syntax
 - Unification algorithm
- 2 Generalisation to binding signatures
- 3 Categorical generalisation
 - A case study: syntax of pure λ -calculus
 - Generic pattern unification
- 4 Example: System F
- 5 Summary of the generic unification algorithm

Pure λ -calculus as a functor

$\mathbb{F}_m :=$ category of finite cardinals and injections between them

Pure λ -calculus as a functor $\Lambda : \mathbb{F}_m \rightarrow \mathbf{Set}$

$$\Lambda_n = \{t \mid \cdot; n \vdash t\}$$

Pure λ -calculus as a fixpoint

$$\Lambda_n \cong \underbrace{\{0, \dots, n-1\}}_{\text{variables}} + \underbrace{\Lambda_n \times \Lambda_n}_{\text{application}} + \underbrace{\Lambda_{n+1}}_{\text{abstraction}}$$

In fact,

$$\Lambda = \mu X. F(X)$$



Initial algebra of the endofunctor F on $[\mathbb{F}_m, \text{Set}]$

$$F(X)_n = \{0, \dots, n-1\} + X_n \times X_n + X_{n+1}$$

Pure λ -calculus with a metavariable

$$\Lambda(M : m)_n = \{t \mid M : m; n \vdash t\}$$

As a fixpoint:

$$\Lambda(M : m) = \mu X. (\underbrace{F(X)}_{\text{operations / variables}} + \textcolor{blue}{arg}^M)$$

$$\begin{aligned} \textcolor{blue}{arg}^M_n &= \{M\text{-arguments in the variable context } \mathbf{n}\} \\ &= \{\vec{x} \in \{0, \dots, n-1\}^m \mid x_1, \dots, x_m \text{ distinct}\} \\ &= \text{hom}_{\mathbb{F}_m}(m, n) \end{aligned}$$

$$\Lambda(M : m) = \mu X. (F(X) + ym)$$

Pure λ -calculus with metavariables

$$\Lambda(\Gamma)_n = \{t \mid \Gamma; n \vdash t\}$$

As a fixpoint:

$$\Lambda(\Gamma) = \mu X. (F(X) + \underbrace{\coprod_{(M:m) \in \Gamma} ym}_{\underline{\Gamma}})$$

$$= \underbrace{T}_{\text{free monad generated by } F}(\underline{\Gamma})$$

$$T(\underline{\Gamma})_n = \{t \mid \Gamma; n \vdash t\}$$

Unification as a Kleisli coequaliser

Claims:

- $\text{hom}(yn, T\underline{\Gamma}) = \text{set of terms in context } \Gamma; n$
- $\text{hom}(\underline{\Gamma}, T\underline{\Delta}) = \text{set of metavariable substitutions } \Gamma \rightarrow \Delta$.
- Most general unifier of $t, u = \text{coequaliser of } yn \xRightarrow[t]{t} T\underline{\Gamma}$

in $\text{Th}_F = \text{Kleisli category restricted to finite coproducts of representable functors } (\underline{\Gamma})$.

Outline

- 1 Pattern unification for pure λ -calculus
 - Syntax
 - Unification algorithm
- 2 Generalisation to binding signatures
- 3 Categorical generalisation
 - A case study: syntax of pure λ -calculus
 - Generic pattern unification
- 4 Example: System F
- 5 Summary of the generic unification algorithm

Signature

- ① \mathcal{A} small category (e.g., \mathbb{F}_m) s.t.
 - all morphisms are monic (pattern restriction)
 - \mathcal{A} has finite connected limits ($M(\vec{x}) \stackrel{?}{=} N(\vec{y})$)
- ② F endofunctor on $[\mathcal{A}, \text{Set}]$ of the shape

$$F(X)_a = \coprod_{o \in O_a} X_{L_{o,1}} \times \cdots \times X_{L_{o,n_o}}$$

such that F restricts to an endofunctor on functors preserving finite connected limits.

Semantics of unification

Claim: Given a signature (\mathcal{A}, F) , a coequaliser diagram in Th_F has a coequaliser as soon as there exists a cocone (i.e., a ‘unifier’)

Proof: By proving termination of the unification algorithm.

Outline

- 1 Pattern unification for pure λ -calculus
 - Syntax
 - Unification algorithm
- 2 Generalisation to binding signatures
- 3 Categorical generalisation
 - A case study: syntax of pure λ -calculus
 - Generic pattern unification
- 4 Example: System F
- 5 Summary of the generic unification algorithm

Example: System F

- Objects of \mathcal{A} :

$$\underbrace{n \mid \tau_1, \dots, \tau_p \vdash \tau_f}_{\text{Types taking free variables in } \{0, \dots, n-1\}}$$

\Rightarrow System F defines a functor $\Phi : \mathcal{A} \rightarrow \text{Set}$

$$\Phi(n \mid \tau_1, \dots, \tau_p \vdash \tau_f) = \{t \text{ s.t. } n \mid \tau_1, \dots, \tau_p \vdash t : \tau_f\}$$

- Endofunctor F on $[\mathcal{A}, \text{Set}]$ s.t. $\Phi = \mu X.F(X)$?
- Metavariables?

The endofunctor for system F

$$F(X)_a = \coprod_{o \in O_a} X_{L_{o,1}} \times \cdots \times X_{L_{o,n_o}}$$

Typing rule	$F(X)_{n \Gamma \vdash \tau} = \coprod \dots$
$\frac{n+1 wk(\Gamma) \vdash t : \tau'}{n \Gamma \vdash \Lambda t : \Lambda \tau'} \text{T-ABS}$	$\coprod_{\tau' \text{ s.t. } \tau = \Lambda \tau'} X_{n+1 wk(\Gamma) \vdash \tau'}$
$\frac{n \Gamma \vdash t : \Lambda \tau_1}{n \Gamma \vdash t \cdot \tau_2 : \tau_1[\tau_2]} \text{T-APP}$	$\coprod_{\tau_1, \tau_2 \text{ s.t. } \tau = \tau_1[\tau_2]} X_{n \Gamma \vdash \Lambda \tau_1}$
$\frac{n \Gamma \vdash t : \tau' \Rightarrow \tau \quad n \Gamma \vdash u : \tau'}{n \Gamma \vdash t \ u : \tau} \text{APP}$	$\coprod_{\tau'} X_{n \Gamma \vdash \tau' \Rightarrow \tau} \times X_{n \Gamma \vdash \tau'}$
\dots	\dots

- System F = initial algebra for F

Metavariables in system F

$$\frac{\alpha_1, \dots, \alpha_p \text{ distinct, } < n \quad x_1, \dots, x_q \text{ distinct} \quad \tau_i[\vec{\alpha}] = \Gamma_i}{\underbrace{\Delta, M : (p|\tau_1, \dots, \tau_q \vdash \tau_f); n|\Gamma \vdash M(\underbrace{\vec{\alpha}}_{\text{type variables}}, \vec{x}) : \tau_f[\vec{\alpha}]}_{\text{Metavariable context}}}$$

Unification in system F

$$M(\vec{\alpha}, \vec{x}) \stackrel{?}{=} M(\vec{\beta}, \vec{y})$$

Most general unifier: $M \mapsto N(\vec{\gamma}, \vec{z})$, where

- $\vec{\gamma}$ maximal s.t.

$$\alpha_{\vec{\gamma}} = \beta_{\vec{\gamma}}$$

- \vec{z} maximal s.t.

$$x'_{\vec{z}} = y'_{\vec{z}}$$

where $\vec{x}' = \vec{x}[\vec{\gamma}^{-1}]$ and $\vec{y}' = \vec{y}[\vec{\gamma}^{-1}]$

Outline

- 1 Pattern unification for pure λ -calculus
 - Syntax
 - Unification algorithm
- 2 Generalisation to binding signatures
- 3 Categorical generalisation
 - A case study: syntax of pure λ -calculus
 - Generic pattern unification
- 4 Example: System F
- 5 Summary of the generic unification algorithm

Interpreting the unification statements

Notations

$$\begin{array}{ll} \Gamma \vdash t = u \Rightarrow \sigma \vdash \Delta & \Leftrightarrow \quad \sigma = \text{coequaliser of } t \text{ and } u \\ \Gamma \vdash t :> f \Rightarrow v; \sigma \vdash \Delta & \Leftrightarrow \quad \sigma, v = \text{pushout of } t \text{ and } f \end{array}$$

mostly used in $\text{Th}_F^* = \text{Th}_F + \text{a free terminal object } \perp$.

Summary of the generic unification algorithm

$$\begin{array}{c}
 \overline{\Gamma \vdash () = () \Rightarrow 1_\Gamma \dashv \Gamma} \quad \overline{\perp \vdash \vec{t} = \vec{u} \Rightarrow ! \dashv \perp} \\
 \\
 \frac{\Gamma \vdash t_1 = u_1 \Rightarrow \sigma_1 \dashv \Delta_1 \quad \Delta_1 \vdash \vec{t}_2[\sigma_1] = \vec{u}_2[\sigma_1] \Rightarrow \sigma_2 \dashv \Delta_2}{\Gamma \vdash t_1, \vec{t}_2 = u_1, \vec{u}_2 \Rightarrow \sigma_1[\sigma_2] \dashv \Delta_2} \text{U-SPLIT} \\
 \\
 \frac{\Gamma \vdash \vec{t} = \vec{u} \Rightarrow \sigma \dashv \Delta}{\Gamma \vdash o(\vec{t}) = o(\vec{u}) \Rightarrow \sigma \dashv \Delta} \text{U-RIGRIG} \quad \frac{o \neq o'}{\Gamma \vdash o(\vec{t}) = o'(\vec{u}) \Rightarrow ! \dashv \perp} \\
 \\
 \frac{u|_\Gamma = u' \quad \Gamma \vdash u' :> M(x) \Rightarrow v; \sigma \dashv \Delta}{\Gamma, M : b \vdash M(x) = u \Rightarrow \sigma, M \mapsto v \dashv \Delta} \text{U-NOCYCLE} + \text{sym} \\
 \\
 \frac{b \vdash x =_{\mathcal{D}} y \Rightarrow z \dashv c}{\Gamma, M : b \vdash M(x) = M(y) \Rightarrow M \mapsto M'(z) \dashv \Gamma, M' : c} \text{U-FLEXFLEX} \\
 \\
 \frac{u = o(\vec{t}) \quad u|_\Gamma \neq \dots}{\Gamma, M : b \vdash M(x) = u \Rightarrow ! \dashv \perp} \text{U-CYCLIC} + \text{sym}
 \end{array}$$

Pruning phase

$$\overline{\Gamma \vdash () :> () \Rightarrow (); 1_\Gamma \dashv \Gamma} \quad \overline{\perp \vdash \vec{t} :> \vec{f} \Rightarrow !; ! \dashv \perp}$$

$$\frac{\Gamma \vdash t_1 :> f_1 \Rightarrow u_1; \sigma_1 \dashv \Delta_1 \quad \Delta_1 \vdash \vec{t}_2[\sigma_1] :> \vec{f}_2 \Rightarrow \vec{u}_2; \sigma_2 \dashv \Delta_2}{\Gamma \vdash t_1, \vec{t}_2 :> f_1 + \vec{f}_2 \Rightarrow u_1[\sigma_2], \vec{u}_2; \sigma_1[\sigma_2] \dashv \Delta_2} \text{P-SPLIT}$$

$$\frac{\Gamma \vdash \vec{t} :> \mathcal{L}^+ x^o \Rightarrow \vec{u}; \sigma \dashv \Delta \quad o = x \cdot o'}{\Gamma \vdash o(\vec{t}) :> N(x) \Rightarrow o'(\vec{u}); \sigma \dashv \Delta} \text{P-RIG}$$

$$\frac{o \neq x \cdot \dots}{\Gamma \vdash o(\vec{t}) :> N(x) \Rightarrow !; ! \dashv \perp}$$

$$\frac{c \vdash_{\mathcal{D}} y :> x \Rightarrow y'; x' \dashv d}{\Gamma, M : c \vdash M(y) :> N(x) \Rightarrow M'(y'); M \mapsto M'(x') \dashv \Gamma, M' : d} \text{P-FLEX}$$